

```
In [ ]: !pip install catboost
```

```
Requirement already satisfied: catboost in /usr/local/lib/python3.11/dist-packages (1.2.8)
Requirement already satisfied: graphviz in /usr/local/lib/python3.11/dist-packages (from catboost) (0.20.3)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (from catboost) (3.10.0)
Requirement already satisfied: numpy<3.0,>=1.16.0 in /usr/local/lib/python3.11/dist-packages (from catboost) (2.0.2)
Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.11/dist-packages (from catboost) (2.2.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from catboost) (1.14.1)
Requirement already satisfied: plotly in /usr/local/lib/python3.11/dist-packages (from catboost) (5.24.1)
Requirement already satisfied: six in /usr/local/lib/python3.11/dist-packages (from catboost) (1.17.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.24->catboost) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (4.57.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (1.4.8)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (24.2)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (11.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (3.2.3)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.11/dist-packages (from plotly->catboost) (9.1.2)
```

```
In [ ]: import numpy as np
import pandas as pd
from functools import reduce
import matplotlib.pyplot as plt
import seaborn as sns
import time
from itertools import combinations

from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split, KFold
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.neighbors import KNeighborsRegressor
```

```

from sklearn.metrics import mean_squared_error, r2_score, mean_abso

import xgboost as xgb
from xgboost import XGBRegressor
from catboost import CatBoostRegressor
import lightgbm as lgb
from lightgbm import LGBMRegressor
import torch

```

```

In [ ]: # Load the file with VISDATE as well
csv_file_path3 = 'GENETIC_14Apr2025.csv'
full_df = pd.read_csv(csv_file_path3)
# Updated list of column names, focusing on volumetric measurements
selected_columns = [
    'RID', 'APVOLUME'
]
# Create a dictionary of DataFrames, each containing only one of th
dataframe3 = pd.read_csv(csv_file_path3, usecols=selected_columns)
# This will create DataFrames for the columns that exist in the CSV
# preventing KeyError in case some columns are missing
print(dataframe3.head())

```

	RID	APVOLUME
0	2	10.0
1	3	10.0
2	4	8.0
3	5	10.0
4	7	9.0

```

In [ ]: csv_file_path4 = '/content/UPENNBIOIMK_ROCHE_ELECSYS_14Apr2025.csv'

full_df = pd.read_csv(csv_file_path4)
# Updated list of column names, focusing on volumetric measurements
selected_columns = [
    'RID', 'ABETA42', 'TAU', 'PTAU'
]
# Create a dictionary of DataFrames, each containing only one of th
dataframe4 = pd.read_csv(csv_file_path4, usecols=selected_columns)
# This will create DataFrames for the columns that exist in the CSV
# preventing KeyError in case some columns are missing
print(dataframe4.head())

```

	RID	ABETA42	TAU	PTAU
0	3	741.5	239.7	22.83
1	3	601.4	251.7	24.18
2	4	1501.0	153.1	13.29
3	4	1176.0	159.7	13.30
4	5	547.3	337.0	33.43

```

In [ ]: import pandas as pd

csv_file_path5 = 'UGOTPTAU181_06_18_20_14Apr2025 (1).csv'

full_df = pd.read_csv(csv_file_path4)
# Updated list of column names, focusing on volumetric measurements
selected_columns = [
    'RID', 'PLASMAPTAU181',

```

```

]
# Create a dictionary of DataFrames, each containing only one of th
dataframe5 = pd.read_csv(csv_file_path5, usecols=selected_columns)
# This will create DataFrames for the columns that exist in the CSV
# preventing KeyError in case some columns are missing
print(dataframe5.head())

```

	RID	PLASMAPTAU181
0	2	11.939
1	2	12.936
2	2	13.563
3	2	15.506
4	8	18.305

```

In [ ]: csv_file_path = 'UCSFFSX7_14Apr2025.csv'

full_df = pd.read_csv(csv_file_path)
# Updated list of column names, focusing on volumetric measurements
"""selected_columns = [
"ST58CV", "ST58SA", "ST58TA", "ST58TS", # Left Superior Temporal
"ST117CV", "ST117SA", "ST117TA", "ST117TS", # Right Superior Temporal
"ST40CV", "ST40SA", "ST40TA", "ST40TS", # Left Middle Temporal
"ST99CV", "ST99SA", "ST99TA", "ST99TS", # Right Middle Temporal
"ST32CV", "ST32SA", "ST32TA", "ST32TS", # Left Inferior Temporal
"ST91CV", "ST91SA", "ST91TA", "ST91TS", # Right Inferior Temporal
"ST60CV", "ST60SA", "ST60TA", "ST60TS", # Left Temporal Pole
"ST119CV", "ST119SA", "ST119TA", "ST119TS", # Right Temporal Pole
"ST62CV", "ST62SA", "ST62TA", "ST62TS", # Left Transverse Temporal
"ST121CV", "ST121SA", "ST121TA", "ST121TS" # Right Transverse Temporal
]"""
selected_columns = [
"RID",
"ST58TA", # Cortical Thickness Average of Left Superior Temporal
"ST117TA", # Cortical Thickness Average of Right Superior Temporal
"ST40TA", # Cortical Thickness Average of Left Middle Temporal
"ST99TA", # Cortical Thickness Average of Right Middle Temporal
"ST32TA", # Cortical Thickness Average of Left Inferior Temporal
"ST91TA", # Cortical Thickness Average of Right Inferior Temporal
"ST60TA", # Cortical Thickness Average of Left Temporal Pole
"ST119TA", # Cortical Thickness Average of Right Temporal Pole
"ST62TA", # Cortical Thickness Average of Left Transverse Temporal
"ST121TA" # Cortical Thickness Average of Right Transverse Temporal
]
# Create a dictionary of DataFrames, each containing only one of th
dataframe9 = pd.read_csv(csv_file_path, usecols=selected_columns)
# This will create DataFrames for the columns that exist in the CSV
# preventing KeyError in case some columns are missing
print(dataframe9.head())

```

	RID	ST117TA	ST119TA	ST121TA	ST32TA	ST40TA	ST58TA	ST60TA
ST62TA \								
0	4213	2.349	3.802	2.145	2.703	2.568	2.471	3.568
2.095								
1	4453	2.571	3.739	2.360	2.509	2.560	2.596	3.785
2.348								
2	4104	2.912	3.865	2.311	2.706	2.663	2.756	3.508
2.421								
3	2153	2.954	3.855	2.770	2.758	2.780	2.922	3.402
2.448								
4	4303	2.554	3.347	2.266	2.693	2.754	2.644	3.303
2.233								

	ST91TA	ST99TA
0	2.463	2.498
1	2.734	2.651
2	2.862	2.903
3	2.641	2.801
4	2.824	2.774

```
In [ ]: csv_file_path = 'PTDEM0G_14Apr2025 (1).csv'

full_df = pd.read_csv(csv_file_path)
# Updated list of column names, focusing on volumetric measurements
selected_columns = [
    "RID",
    # Demographic and Background Information
    "PTGENDER", "PTDOB",
    # Clinical and Site-Specific Data
    "PTDOBY", "PTMARRY", "PTEDUCAT", "PTETHCAT",
]
# Create a dictionary of DataFrames, each containing only one of the
dataframe7 = pd.read_csv(csv_file_path, usecols=selected_columns)
# This will create DataFrames for the columns that exist in the CSV
# preventing KeyError in case some columns are missing
print(dataframe7.head())
```

	RID	PTGENDER	PTDOB
0	2	1.0	04/1931
1	1	2.0	12/1944
2	3	1.0	05/1924
3	4	1.0	01/1938
4	5	1.0	12/1931

```
In [ ]: csv_file_path = 'MEDHIST_14Apr2025.csv'

full_df = pd.read_csv(csv_file_path)
# Updated list of column names, focusing on volumetric measurements
selected_columns = [
    "RID",
    # Medical Conditions Related to Brain Structure
    "MH14ALCH", # Alcohol Abuse
    "MH15DRUG", # Drug Abuse
    "MH16SMOK", # Smoking
    "MH2NEURL", # Neurologic (other than AD)
    "MHPSYCH" # Psychiatric Conditions
]
```

```
# Create a dictionary of DataFrames, each containing only one of the
dataframe8 = pd.read_csv(csv_file_path, usecols=selected_columns)
# This will create DataFrames for the columns that exist in the CSV
# preventing KeyError in case some columns are missing
print(dataframe8.head())
```

	RID	MHPSYCH	MH2NEURL	MH14ALCH	MH15DRUG	MH16SMOK
0	2	0	0	0	0	0
1	1	0	0	0	0	0
2	3	0	0	0	0	1
3	4	0	0	0	0	1
4	5	0	0	0	0	1

```
In [ ]: dfs = [dataframe3, dataframe4, dataframe5, dataframe7, dataframe8,

# Check if 'RID' is present in each DataFrame
rid_check = [True if 'RID' in df.columns else False for df in dfs]

# Display all columns when printing DataFrames
pd.set_option('display.max_columns', None)

# Print the check results
print("RID presence in each DataFrame:", rid_check)

# Merge all DataFrames on 'RID' using inner join
merged_df = reduce(lambda left, right: pd.merge(left, right, on='RI

# Display the merged DataFrame
print(merged_df.head())
```

RID presence in each DataFrame: [True, True, True, True, True, True]

RID	APVOLUME	ABETA42	TAU	PTAU	PLASMAPTAU181	PTGENDER
PTDOB \						
0 31	10.0	1774.0	266.8	22.55	18.642	2.0 0
1/1928						
1 31	10.0	1774.0	266.8	22.55	18.642	2.0 0
1/1928						
2 31	10.0	1774.0	266.8	22.55	18.642	2.0 0
1/1928						
3 31	10.0	1774.0	266.8	22.55	18.642	2.0 0
1/1928						
4 31	10.0	1774.0	266.8	22.55	18.642	2.0 0
1/1928						

MHPSYCH	MH2NEURL	MH14ALCH	MH15DRUG	MH16SMOK	ST117TA	ST119TA
ST121TA \						
0 0	0	0	0	0	2.318	3.495
2.080						
1 0	0	0	0	0	2.412	2.758
2.469						
2 0	0	0	0	0	2.318	3.495
2.080						
3 0	0	0	0	0	2.412	2.758
2.469						
4 0	0	0	0	0	2.318	3.495
2.080						

ST32TA	ST40TA	ST58TA	ST60TA	ST62TA	ST91TA	ST99TA
0 2.767	2.740	2.304	3.670	1.718	2.528	2.463
1 2.744	2.738	2.548	3.687	2.177	2.229	2.384
2 2.767	2.740	2.304	3.670	1.718	2.528	2.463
3 2.744	2.738	2.548	3.687	2.177	2.229	2.384
4 2.767	2.740	2.304	3.670	1.718	2.528	2.463

```
In [ ]: pd.set_option('display.max_columns', None)
print(merged_df.head())
```

	RID	APVOLUME	ABETA42	TAU	PTAU	PLASMAPTAU181	PTGENDER
PTDOB \							
0	31	10.0	1774.0	266.8	22.55	18.642	2.0 0
1/1928							
1	31	10.0	1774.0	266.8	22.55	18.642	2.0 0
1/1928							
2	31	10.0	1774.0	266.8	22.55	18.642	2.0 0
1/1928							
3	31	10.0	1774.0	266.8	22.55	18.642	2.0 0
1/1928							
4	31	10.0	1774.0	266.8	22.55	18.642	2.0 0
1/1928							

	MHPSYCH	MH2NEURL	MH14ALCH	MH15DRUG	MH16SMOK	ST117TA	ST119TA
ST121TA \							
0	0	0	0	0	0	2.318	3.495
2.080							
1	0	0	0	0	0	2.412	2.758
2.469							
2	0	0	0	0	0	2.318	3.495
2.080							
3	0	0	0	0	0	2.412	2.758
2.469							
4	0	0	0	0	0	2.318	3.495
2.080							

	ST32TA	ST40TA	ST58TA	ST60TA	ST62TA	ST91TA	ST99TA
0	2.767	2.740	2.304	3.670	1.718	2.528	2.463
1	2.744	2.738	2.548	3.687	2.177	2.229	2.384
2	2.767	2.740	2.304	3.670	1.718	2.528	2.463
3	2.744	2.738	2.548	3.687	2.177	2.229	2.384
4	2.767	2.740	2.304	3.670	1.718	2.528	2.463

```
In [ ]: print(merged_df.describe())
```

	RID	APVOLUME	ABETA42	TAU \
count	742688.000000	579314.000000	740998.000000	740588.000000
mean	2545.892749	0.444097	1238.469569	271.269763
std	1516.797451	6.235802	637.193384	132.301563
min	23.000000	-4.000000	203.000000	80.080000
25%	1261.000000	-4.000000	702.400000	186.400000
50%	2245.000000	-4.000000	1106.000000	238.800000
75%	4175.000000	9.000000	1689.000000	310.000000
max	5296.000000	12.000000	3949.000000	1018.000000

	PTAU	PLASMAPTAU181	PTGENDER	MHPSYCH \
count	737202.000000	742688.000000	703651.000000	742688.000000
mean	25.252673	16.011811	1.498388	0.405950
std	14.873441	9.512347	0.640717	0.491075
min	8.260000	0.468000	-4.000000	0.000000
25%	16.300000	10.300000	1.000000	0.000000
50%	21.090000	14.564000	2.000000	0.000000
75%	28.580000	19.264000	2.000000	1.000000
max	103.700000	451.398000	2.000000	1.000000

	MH2NEURL	MH14ALCH	MH15DRUG	MH16SMOK \
--	----------	----------	----------	------------

count	742688.000000	742688.000000	742688.000000	742688.000000
mean	0.386790	0.029850	0.011048	0.402071
std	0.487015	0.170173	0.104526	0.490316
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000

	ST117TA	ST119TA	ST121TA	ST32TA \
count	736958.000000	736958.000000	742128.000000	736958.000000
mean	2.591584	3.453158	2.305323	2.645257
std	0.197361	0.374494	0.218933	0.181849
min	1.644000	1.407000	1.413000	1.729000
25%	2.469000	3.295000	2.159000	2.540000
50%	2.601000	3.517000	2.310000	2.653000
75%	2.710000	3.719750	2.450000	2.767000
max	3.124000	4.315000	2.970000	3.197000

	ST40TA	ST58TA	ST60TA	ST62TA \
count	736958.000000	736958.000000	736958.000000	742128.000000
mean	2.659840	2.566967	3.373313	2.273667
std	0.182294	0.195884	0.372920	0.206554
min	1.773000	1.723000	1.423000	1.396000
25%	2.563000	2.457000	3.190000	2.148000
50%	2.673000	2.583000	3.424000	2.278000
75%	2.782000	2.702000	3.606000	2.413000
max	3.147000	3.010000	4.310000	2.884000

	ST91TA	ST99TA
count	736958.000000	736958.000000
mean	2.669802	2.692279
std	0.176935	0.180707
min	1.747000	1.661000
25%	2.573000	2.598000
50%	2.694000	2.703000
75%	2.805000	2.802000
max	3.195000	3.168000

```
In [ ]: print(merged_df.isnull().sum())
```



```


RID                0
APVOLUME           163374
ABETA42            1690
TAU                2100
PTAU              5486
PLASMAPTAU181      0
PTGENDER           39037
PTDOB             42802
MHPSYCH            0
MH2NEURL           0
MH14ALCH           0
MH15DRUG           0
MH16SMOK           0
ST117TA           5730
ST119TA           5730
ST121TA           560
ST32TA            5730
ST40TA            5730
ST58TA            5730
ST60TA            5730
ST62TA            560
ST91TA            5730
ST99TA            5730
dtype: int64

```

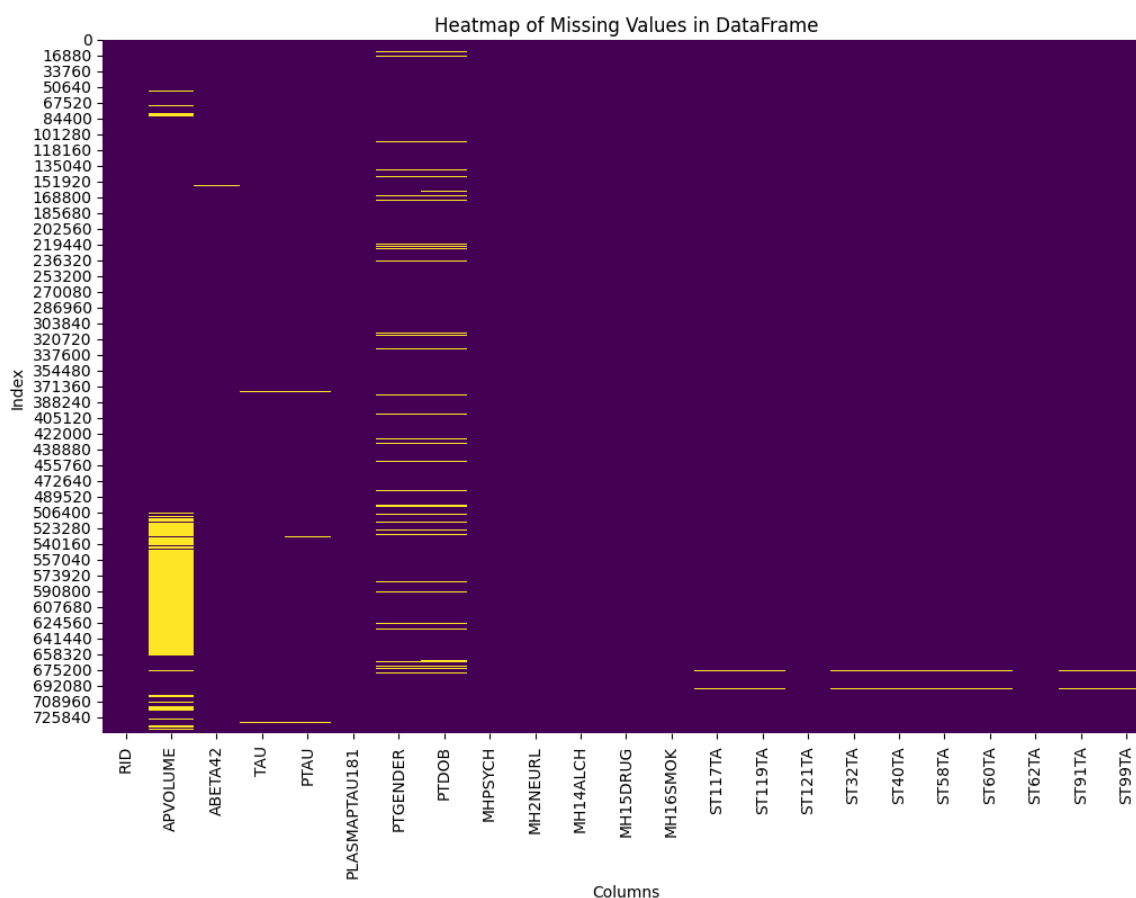
```

In [ ]: missing_values = merged_df.isnull().sum()
missing_values.sort_values(inplace=True) # Sort the values
missing_values.plot(kind='bar', figsize=(10, 6)) # Create a bar plot
plt.xlabel('Columns')
plt.ylabel('Number of Missing Values')
plt.title('Number of Missing Values in Each Column')
plt.show()

```



```
In [ ]: plt.figure(figsize=(12, 8))
sns.heatmap(merged_df.isnull(), cbar=False, cmap='viridis')
plt.title('Heatmap of Missing Values in DataFrame')
plt.xlabel('Columns')
plt.ylabel('Index')
plt.show()
```



```
In [ ]: cleaned_df_dropna = merged_df.dropna()
print(cleaned_df_dropna.describe())
```

	RID	APVOLUME	ABETA42	TAU \
count	537865.000000	537865.000000	537865.000000	537865.000000
mean	2530.891097	0.390705	1261.005796	266.261582
std	1533.006742	6.215089	643.410172	115.849659
min	23.000000	-4.000000	203.000000	88.690000
25%	1261.000000	-4.000000	717.200000	186.600000
50%	2245.000000	-4.000000	1146.000000	239.100000
75%	4187.000000	9.000000	1718.000000	313.200000
max	5296.000000	12.000000	3949.000000	915.800000

	PTAU	PLASMAPTAU181	PTGENDER	MHPSYCH \
count	537865.000000	537865.000000	537865.000000	537865.000000
mean	24.608651	16.228029	1.512430	0.400106
std	12.965423	9.808985	0.499846	0.489920
min	8.260000	0.468000	1.000000	0.000000
25%	16.400000	10.274000	1.000000	0.000000
50%	21.090000	14.533000	2.000000	0.000000
75%	28.710000	19.988000	2.000000	1.000000
max	103.700000	451.398000	2.000000	1.000000

	MH2NEURL	MH14ALCH	MH15DRUG	MH16SMOK \
count	537865.000000	537865.000000	537865.000000	537865.000000
mean	0.378498	0.028171	0.012076	0.425254
std	0.485013	0.165460	0.109223	0.494382
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000

	ST117TA	ST119TA	ST121TA	ST32TA \
count	537865.000000	537865.000000	537865.000000	537865.000000
mean	2.582546	3.431361	2.303026	2.637437
std	0.203246	0.384000	0.216833	0.185909
min	1.644000	1.407000	1.413000	1.729000
25%	2.457000	3.274000	2.161000	2.532000
50%	2.593000	3.503000	2.309000	2.648000
75%	2.704000	3.692000	2.443000	2.762000
max	3.124000	4.315000	2.970000	3.197000

	ST40TA	ST58TA	ST60TA	ST62TA \
count	537865.000000	537865.000000	537865.000000	537865.000000
mean	2.652819	2.560486	3.347277	2.270687
std	0.188352	0.201863	0.378968	0.211328
min	1.773000	1.723000	1.423000	1.396000
25%	2.553000	2.446000	3.164000	2.141000
50%	2.667000	2.578000	3.411000	2.275000
75%	2.780000	2.701000	3.581000	2.413000
max	3.147000	3.010000	4.310000	2.884000

	ST91TA	ST99TA
count	537865.000000	537865.000000
mean	2.663668	2.680271
std	0.182223	0.182187
min	1.747000	1.661000
25%	2.565000	2.586000
50%	2.688000	2.692000
75%	2.800000	2.796000
max	3.195000	3.168000

```
In [ ]: cleaned_df_mean = merged_df.copy()

for column in cleaned_df_mean.select_dtypes(include=['float64', 'in
cleaned_df_mean[column].fillna(cleaned_df_mean[column].mean(),
```

<ipython-input-19-eafc50d4d64f>:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
cleaned_df_mean[column].fillna(cleaned_df_mean[column].mean(), inplace=True)
```

```
In [ ]: def parse_date(date_str):
        try:
            # Try parsing as "4/1/1931" format
            return pd.to_datetime(date_str, format='%m/%d/%Y')
        except ValueError:
            try:
                # Try parsing as "Apr-31" format
                month, year = date_str.split('-')
                month_num = pd.to_datetime(month, format='%b').month
                return pd.to_datetime(f'1-{month_num}-{year}', format='%d-%m-%Y')
            except ValueError:
                try:
                    # Try parsing as "01/1934" format
                    month, year = date_str.split('/')
                    return pd.to_datetime(f'1-{month}-{year}', format='%d-%m-%Y')
                except ValueError:
                    # Return NaT for anything else
                    return pd.NaT

        # Apply it
        cleaned_df_dropna['PTDOB'] = cleaned_df_dropna['PTDOB'].apply(parse_date)

        # Print a few values to confirm
        print(cleaned_df_dropna['PTDOB'].head())
```

```
0    1928-01-01
1    1928-01-01
2    1928-01-01
3    1928-01-01
4    1928-01-01
Name: PTDOB, dtype: datetime64[ns]
```

<ipython-input-20-a20a3a60c4fb>:21: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
cleaned_df_dropna['PTDOB'] = cleaned_df_dropna['PTDOB'].apply(parse_date)
```

```
In [ ]: target_names = [
    "ST58TA", # Cortical Thickness Average of Left Superior Temporal
    "ST117TA", # Cortical Thickness Average of Right Superior Temporal
    "ST40TA", # Cortical Thickness Average of Left Middle Temporal
    "ST99TA", # Cortical Thickness Average of Right Middle Temporal
    "ST32TA", # Cortical Thickness Average of Left Inferior Temporal
    "ST91TA", # Cortical Thickness Average of Right Inferior Temporal
    "ST60TA", # Cortical Thickness Average of Left Temporal Pole
    "ST119TA", # Cortical Thickness Average of Right Temporal Pole
    "ST62TA", # Cortical Thickness Average of Left Transverse Temporal
    "ST121TA" # Cortical Thickness Average of Right Transverse Temporal
]
# Features DataFrame (X)
# Drop target columns from the main DataFrame to create the feature
X = cleaned_df_dropna.drop(columns=target_names + ['RID'])
# Targets DataFrame (y)
# Select only the target columns for the target DataFrame
y = cleaned_df_dropna[target_names]
```

```
In [ ]: X['PTDOB'] = pd.to_datetime(X['PTDOB'])
X['year'] = X['PTDOB'].dt.year
X['month'] = X['PTDOB'].dt.month
X['day'] = X['PTDOB'].dt.day
# Now, you can drop the original datetime column if needed
X = X.drop(columns=['PTDOB'])
```

```
In [ ]: import pandas as pd
from sklearn.preprocessing import MinMaxScaler

# Columns to use for outlier detection and scaling
cols = ['APVOLUME', 'ABETA42', 'TAU', 'PTAU', 'PLASMAPTAU181']

# Calculate Q1, Q3, and IQR
Q1 = X[cols].quantile(0.25)
Q3 = X[cols].quantile(0.75)
IQR = Q3 - Q1

# Define an outlier mask (True = not an outlier)
outlier_mask = ~(X[cols] < (Q1 - 1.5 * IQR)) | (X[cols] > (Q3 + 1.5 * IQR))

# Apply the mask to filter out outliers
X_clean = X.loc[outlier_mask, cols]
y_clean = y[outlier_mask]

# Normalize the data using Min-Max Scaling
scaler = MinMaxScaler()
X_scaled = pd.DataFrame(scaler.fit_transform(X_clean), columns=cols)

# Display shapes before and after cleaning
print("Original shape of X:", X.shape)
print("New shape of X after outlier removal and scaling:", X_scaled.shape)
print("Original shape of y:", y.shape)
print("New shape of y after outlier and scaling process:", y_clean.shape)
```

Original shape of X: (537865, 14)


New shape of X after outlier removal and scaling: (484436, 5)

Original shape of y: (537865, 10)

New shape of y after outlier and scaling process: (484436, 10)

```
In [ ]: correlation_matrix = X.corr(method='spearman')
```

```
In [ ]: plt.figure(figsize=(12, 10)) # Set the figure size as needed
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap='coolwa
cbar=True, square=True, linewidths=.5)
plt.title('Spearman Correlation Heatmap')
plt.show()
```



```
In [ ]: cleaned_df_dropna.shape
```

```
Out[ ]: (537865, 23)
```

```
In [ ]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np

# Split the dataset into training and testing sets
```

```

X_train_full, X_test, y_train_full, y_test = train_test_split(X, y,

# Now split the training set into four parts, simulating distributi
parts_X = np.array_split(X_train_full, 4)
parts_y = np.array_split(y_train_full, 4)

# Define the models to be used
models = {
    'Linear Regression': LinearRegression(),
    'Decision Tree': DecisionTreeRegressor(),
    'Random Forest': RandomForestRegressor(n_estimators=100),
    'k-NN': KNeighborsRegressor(n_neighbors=10)
}

# Train models on each client's data for each target and store pred
client_predictions = {
    name: {col: [] for col in y_train_full.columns}
    for name in models.keys()
}

for i in range(4):
    X_train, y_train = parts_X[i], parts_y[i]
    for name, model in models.items():
        for col in y_train.columns:
            model.fit(X_train, y_train[col])
            predictions = model.predict(X_test)
            client_predictions[name][col].append(predictions)

# Average predictions from each client for each target variable and
for name, targets in client_predictions.items():
    print(f"\nResults for {name}:")
    for target_col, predictions in targets.items():
        average_prediction = np.mean(predictions, axis=0)
        mse = mean_squared_error(y_test[target_col], average_prediction)
        r2 = r2_score(y_test[target_col], average_prediction)
        print(f" {target_col} - Averaged MSE: {mse:.4f}, Averaged

```

```

/usr/local/lib/python3.11/dist-packages/numpy/_core/fromnumeric.py:5
7: FutureWarning: 'DataFrame.swapaxes' is deprecated and will be rem
oved in a future version. Please use 'DataFrame.transpose' instead.
    return bound(*args, **kwargs)

```

Results for Linear Regression:

ST58TA – Averaged MSE: 0.0287, Averaged R² Score: 0.2937
 ST117TA – Averaged MSE: 0.0287, Averaged R² Score: 0.3011
 ST40TA – Averaged MSE: 0.0287, Averaged R² Score: 0.1904
 ST99TA – Averaged MSE: 0.0277, Averaged R² Score: 0.1649
 ST32TA – Averaged MSE: 0.0302, Averaged R² Score: 0.1293
 ST91TA – Averaged MSE: 0.0282, Averaged R² Score: 0.1500
 ST60TA – Averaged MSE: 0.1342, Averaged R² Score: 0.0553
 ST119TA – Averaged MSE: 0.1285, Averaged R² Score: 0.1191
 ST62TA – Averaged MSE: 0.0389, Averaged R² Score: 0.1289
 ST121TA – Averaged MSE: 0.0398, Averaged R² Score: 0.1548

Results for Decision Tree:

ST58TA – Averaged MSE: 0.0042, Averaged R² Score: 0.8959
 ST117TA – Averaged MSE: 0.0038, Averaged R² Score: 0.9087
 ST40TA – Averaged MSE: 0.0043, Averaged R² Score: 0.8793
 ST99TA – Averaged MSE: 0.0039, Averaged R² Score: 0.8820
 ST32TA – Averaged MSE: 0.0053, Averaged R² Score: 0.8474
 ST91TA – Averaged MSE: 0.0059, Averaged R² Score: 0.8217
 ST60TA – Averaged MSE: 0.0276, Averaged R² Score: 0.8055
 ST119TA – Averaged MSE: 0.0324, Averaged R² Score: 0.7778
 ST62TA – Averaged MSE: 0.0067, Averaged R² Score: 0.8497
 ST121TA – Averaged MSE: 0.0081, Averaged R² Score: 0.8288

Results for Random Forest:

ST58TA – Averaged MSE: 0.0042, Averaged R² Score: 0.8969
 ST117TA – Averaged MSE: 0.0037, Averaged R² Score: 0.9096
 ST40TA – Averaged MSE: 0.0042, Averaged R² Score: 0.8805
 ST99TA – Averaged MSE: 0.0039, Averaged R² Score: 0.8831
 ST32TA – Averaged MSE: 0.0052, Averaged R² Score: 0.8491
 ST91TA – Averaged MSE: 0.0059, Averaged R² Score: 0.8235
 ST60TA – Averaged MSE: 0.0273, Averaged R² Score: 0.8078
 ST119TA – Averaged MSE: 0.0321, Averaged R² Score: 0.7799
 ST62TA – Averaged MSE: 0.0067, Averaged R² Score: 0.8512
 ST121TA – Averaged MSE: 0.0080, Averaged R² Score: 0.8304

Results for k-NN:

ST58TA – Averaged MSE: 0.0045, Averaged R² Score: 0.8884
 ST117TA – Averaged MSE: 0.0040, Averaged R² Score: 0.9021
 ST40TA – Averaged MSE: 0.0045, Averaged R² Score: 0.8724
 ST99TA – Averaged MSE: 0.0041, Averaged R² Score: 0.8758
 ST32TA – Averaged MSE: 0.0055, Averaged R² Score: 0.8413
 ST91TA – Averaged MSE: 0.0061, Averaged R² Score: 0.8162
 ST60TA – Averaged MSE: 0.0284, Averaged R² Score: 0.7998
 ST119TA – Averaged MSE: 0.0335, Averaged R² Score: 0.7707
 ST62TA – Averaged MSE: 0.0070, Averaged R² Score: 0.8427
 ST121TA – Averaged MSE: 0.0084, Averaged R² Score: 0.8216

```

In [ ]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error, r2_score

```



```

# Assuming 'X' and 'y' are defined (features and targets)
# Replace with actual dataset loading if needed
# Example:
# data = pd.read_csv("cleaned_data.csv")
# X = data.drop(columns=['target1', 'target2']) # replace with act
# y = data[['target1', 'target2']]

# Split the entire dataset into a training set and a testing set
X_train_full, X_test, y_train_full, y_test = train_test_split(X, y,

# Now split the training set into four parts, simulating distributi
parts_X = np.array_split(X_train_full, 4)
parts_y = np.array_split(y_train_full, 4)

models = {
    'Linear Regression': LinearRegression(),
    'Decision Tree': DecisionTreeRegressor(),
    'Random Forest': RandomForestRegressor(n_estimators=100),
    'k-NN': KNeighborsRegressor(n_neighbors=10)
}

# Train models on each client's data for each target and store pred
client_predictions = {name: {col: [] for col in y_train_full.column
weights = [len(part) for part in parts_y] # Number of samples in e

for i in range(4):
    X_train, y_train = parts_X[i], parts_y[i]
    for name, model in models.items():
        for col in y_train.columns:
            model.fit(X_train, y_train[col])
            predictions = model.predict(X_test)
            client_predictions[name][col].append(predictions)

# Calculate weighted average of predictions from each client for ea
for name, targets in client_predictions.items():
    print(f"Results for {name}:")
    for target_col, predictions in targets.items():
        # Compute weighted average
        weighted_average_prediction = np.average(predictions, axis=
        mse = mean_squared_error(y_test[target_col], weighted_avera
        r2 = r2_score(y_test[target_col], weighted_average_predicti
        print(f" {target_col} - Weighted Averaged MSE: {mse:.4f}, '

```

```

/usr/local/lib/python3.11/dist-packages/numpy/_core/fromnumeric.py:5
7: FutureWarning: 'DataFrame.swapaxes' is deprecated and will be rem
oved in a future version. Please use 'DataFrame.transpose' instead.
    return bound(*args, **kwds)

```

Results for Linear Regression:

ST58TA - Weighted Averaged MSE: 0.0287, Weighted Averaged R2: 0.2937

ST117TA - Weighted Averaged MSE: 0.0287, Weighted Averaged R2: 0.3011

ST40TA - Weighted Averaged MSE: 0.0287, Weighted Averaged R2: 0.1904

ST99TA - Weighted Averaged MSE: 0.0277, Weighted Averaged R2: 0.16

49
ST32TA – Weighted Averaged MSE: 0.0302, Weighted Averaged R2: 0.12
93
ST91TA – Weighted Averaged MSE: 0.0282, Weighted Averaged R2: 0.15
00
ST60TA – Weighted Averaged MSE: 0.1342, Weighted Averaged R2: 0.05
53
ST119TA – Weighted Averaged MSE: 0.1285, Weighted Averaged R2: 0.1
191
ST62TA – Weighted Averaged MSE: 0.0389, Weighted Averaged R2: 0.12
89
ST121TA – Weighted Averaged MSE: 0.0398, Weighted Averaged R2: 0.1
548
Results for Decision Tree:
ST58TA – Weighted Averaged MSE: 0.0042, Weighted Averaged R2: 0.89
59
ST117TA – Weighted Averaged MSE: 0.0038, Weighted Averaged R2: 0.9
087
ST40TA – Weighted Averaged MSE: 0.0043, Weighted Averaged R2: 0.87
92
ST99TA – Weighted Averaged MSE: 0.0039, Weighted Averaged R2: 0.88
19
ST32TA – Weighted Averaged MSE: 0.0053, Weighted Averaged R2: 0.84
74
ST91TA – Weighted Averaged MSE: 0.0059, Weighted Averaged R2: 0.82
17
ST60TA – Weighted Averaged MSE: 0.0276, Weighted Averaged R2: 0.80
55
ST119TA – Weighted Averaged MSE: 0.0324, Weighted Averaged R2: 0.7
777
ST62TA – Weighted Averaged MSE: 0.0067, Weighted Averaged R2: 0.84
97
ST121TA – Weighted Averaged MSE: 0.0081, Weighted Averaged R2: 0.8
287
Results for Random Forest:
ST58TA – Weighted Averaged MSE: 0.0042, Weighted Averaged R2: 0.89
69
ST117TA – Weighted Averaged MSE: 0.0037, Weighted Averaged R2: 0.9
095
ST40TA – Weighted Averaged MSE: 0.0042, Weighted Averaged R2: 0.88
05
ST99TA – Weighted Averaged MSE: 0.0039, Weighted Averaged R2: 0.88
31
ST32TA – Weighted Averaged MSE: 0.0052, Weighted Averaged R2: 0.84
90
ST91TA – Weighted Averaged MSE: 0.0059, Weighted Averaged R2: 0.82
36
ST60TA – Weighted Averaged MSE: 0.0273, Weighted Averaged R2: 0.80
78
ST119TA – Weighted Averaged MSE: 0.0321, Weighted Averaged R2: 0.7
799
ST62TA – Weighted Averaged MSE: 0.0067, Weighted Averaged R2: 0.85
12
ST121TA – Weighted Averaged MSE: 0.0080, Weighted Averaged R2: 0.8
304
Results for k-NN:

ST58TA – Weighted Averaged MSE: 0.0045, Weighted Averaged R2: 0.88
84
ST117TA – Weighted Averaged MSE: 0.0040, Weighted Averaged R2: 0.9
021
ST40TA – Weighted Averaged MSE: 0.0045, Weighted Averaged R2: 0.87
24
ST99TA – Weighted Averaged MSE: 0.0041, Weighted Averaged R2: 0.87
58
ST32TA – Weighted Averaged MSE: 0.0055, Weighted Averaged R2: 0.84
13
ST91TA – Weighted Averaged MSE: 0.0061, Weighted Averaged R2: 0.81
62
ST60TA – Weighted Averaged MSE: 0.0284, Weighted Averaged R2: 0.79
98
ST119TA – Weighted Averaged MSE: 0.0335, Weighted Averaged R2: 0.7
707
ST62TA – Weighted Averaged MSE: 0.0070, Weighted Averaged R2: 0.84
27
ST121TA – Weighted Averaged MSE: 0.0084, Weighted Averaged R2: 0.8
216

```
In [ ]: !pip install xgboost
```

Requirement already satisfied: xgboost in /usr/local/lib/python3.11/dist-packages (2.1.4)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.0.2)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.21.5)
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from xgboost) (1.14.1)

```
In [ ]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
import xgboost as xgb

# Assuming 'X' and 'y' are defined
# X = features, y = DataFrame with multiple target columns

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Initialize XGBoost Regressor
model = xgb.XGBRegressor(objective='reg:squarederror')

# Train and evaluate model for each target variable
for target_col in y_train.columns:
    print(f"Results for target variable: {target_col}")
    model.fit(X_train, y_train[target_col])
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test[target_col], y_pred)
    r2 = r2_score(y_test[target_col], y_pred)
    print(f"XGBoost – MSE: {mse:.4f}, R^2 Score: {r2:.4f}")
    print()
```

Results for target variable: ST58TA
 XGBoost – MSE: 0.0043, R² Score: 0.8933

Results for target variable: ST117TA
 XGBoost – MSE: 0.0039, R² Score: 0.9057

Results for target variable: ST40TA
 XGBoost – MSE: 0.0044, R² Score: 0.8761

Results for target variable: ST99TA
 XGBoost – MSE: 0.0041, R² Score: 0.8776

Results for target variable: ST32TA
 XGBoost – MSE: 0.0054, R² Score: 0.8450

Results for target variable: ST91TA
 XGBoost – MSE: 0.0060, R² Score: 0.8204

Results for target variable: ST60TA
 XGBoost – MSE: 0.0282, R² Score: 0.8018

Results for target variable: ST119TA
 XGBoost – MSE: 0.0326, R² Score: 0.7769

Results for target variable: ST62TA
 XGBoost – MSE: 0.0069, R² Score: 0.8456

Results for target variable: ST121TA
 XGBoost – MSE: 0.0082, R² Score: 0.8256

```
In [ ]: import pandas as pd
        from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LinearRegression
        from sklearn.tree import DecisionTreeRegressor
        from sklearn.ensemble import RandomForestRegressor
        from sklearn.neighbors import KNeighborsRegressor
        from sklearn.metrics import mean_squared_error, r2_score

        # Assuming X and y are already defined DataFrames
        # X = features, y = target variables (possibly multi-output)

        # Split data into training and testing sets
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

        # Initialize regression models
        models = {
            'Linear Regression': LinearRegression(),
            'Decision Tree': DecisionTreeRegressor(),
            'Random Forest': RandomForestRegressor(n_estimators=100),
            'k-NN': KNeighborsRegressor(n_neighbors=10)
        }

        # Train and evaluate each model for each target variable
        for target_col in y_train.columns:
            print(f"\nResults for target variable: {target_col}")
```

```

for name, model in models.items():
    model.fit(X_train, y_train[target_col])
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test[target_col], y_pred)
    r2 = r2_score(y_test[target_col], y_pred)
    accuracy = model.score(X_test, y_test[target_col]) # Same
    print(f"{name} - MSE: {mse:.4f}, R^2 Score: {r2:.4f}, Accur

```

Results for target variable: ST58TA

Linear Regression - MSE: 0.0287, R^2 Score: 0.2937, Accuracy: 0.2937
 Decision Tree - MSE: 0.0042, R^2 Score: 0.8959, Accuracy: 0.8959
 Random Forest - MSE: 0.0042, R^2 Score: 0.8960, Accuracy: 0.8960
 k-NN - MSE: 0.0045, R^2 Score: 0.8891, Accuracy: 0.8891

Results for target variable: ST117TA

Linear Regression - MSE: 0.0287, R^2 Score: 0.3011, Accuracy: 0.3011
 Decision Tree - MSE: 0.0037, R^2 Score: 0.9090, Accuracy: 0.9090
 Random Forest - MSE: 0.0037, R^2 Score: 0.9090, Accuracy: 0.9090
 k-NN - MSE: 0.0040, R^2 Score: 0.9031, Accuracy: 0.9031

Results for target variable: ST40TA

Linear Regression - MSE: 0.0287, R^2 Score: 0.1904, Accuracy: 0.1904
 Decision Tree - MSE: 0.0043, R^2 Score: 0.8793, Accuracy: 0.8793
 Random Forest - MSE: 0.0043, R^2 Score: 0.8795, Accuracy: 0.8795
 k-NN - MSE: 0.0045, R^2 Score: 0.8716, Accuracy: 0.8716

Results for target variable: ST99TA

Linear Regression - MSE: 0.0277, R^2 Score: 0.1649, Accuracy: 0.1649
 Decision Tree - MSE: 0.0039, R^2 Score: 0.8823, Accuracy: 0.8823
 Random Forest - MSE: 0.0039, R^2 Score: 0.8824, Accuracy: 0.8824
 k-NN - MSE: 0.0042, R^2 Score: 0.8744, Accuracy: 0.8744

Results for target variable: ST32TA

Linear Regression - MSE: 0.0302, R^2 Score: 0.1293, Accuracy: 0.1293
 Decision Tree - MSE: 0.0053, R^2 Score: 0.8476, Accuracy: 0.8476
 Random Forest - MSE: 0.0053, R^2 Score: 0.8478, Accuracy: 0.8478
 k-NN - MSE: 0.0056, R^2 Score: 0.8386, Accuracy: 0.8386

Results for target variable: ST91TA

Linear Regression - MSE: 0.0282, R^2 Score: 0.1500, Accuracy: 0.1500
 Decision Tree - MSE: 0.0059, R^2 Score: 0.8221, Accuracy: 0.8221
 Random Forest - MSE: 0.0059, R^2 Score: 0.8222, Accuracy: 0.8222
 k-NN - MSE: 0.0063, R^2 Score: 0.8104, Accuracy: 0.8104

Results for target variable: ST60TA

Linear Regression - MSE: 0.1342, R^2 Score: 0.0553, Accuracy: 0.0553
 Decision Tree - MSE: 0.0277, R^2 Score: 0.8054, Accuracy: 0.8054
 Random Forest - MSE: 0.0276, R^2 Score: 0.8057, Accuracy: 0.8057
 k-NN - MSE: 0.0293, R^2 Score: 0.7941, Accuracy: 0.7941

Results for target variable: ST119TA

Linear Regression - MSE: 0.1285, R^2 Score: 0.1191, Accuracy: 0.1191
 Decision Tree - MSE: 0.0325, R^2 Score: 0.7773, Accuracy: 0.7773
 Random Forest - MSE: 0.0324, R^2 Score: 0.7776, Accuracy: 0.7776
 k-NN - MSE: 0.0343, R^2 Score: 0.7648, Accuracy: 0.7648

Results for target variable: ST62TA

Linear Regression – MSE: 0.0389, R² Score: 0.1289, Accuracy: 0.1289

Decision Tree – MSE: 0.0067, R² Score: 0.8496, Accuracy: 0.8496

Random Forest – MSE: 0.0067, R² Score: 0.8499, Accuracy: 0.8499

k-NN – MSE: 0.0071, R² Score: 0.8406, Accuracy: 0.8406

Results for target variable: ST121TA

Linear Regression – MSE: 0.0398, R² Score: 0.1548, Accuracy: 0.1548

Decision Tree – MSE: 0.0081, R² Score: 0.8283, Accuracy: 0.8283

Random Forest – MSE: 0.0081, R² Score: 0.8286, Accuracy: 0.8286

k-NN – MSE: 0.0086, R² Score: 0.8179, Accuracy: 0.8179

```
In [ ]: import pandas as pd
import numpy as np
import time
from sklearn.model_selection import train_test_split, KFold
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
import xgboost as xgb
from catboost import CatBoostRegressor
import lightgbm as lgb
import matplotlib.pyplot as plt

# Assuming X and y are already defined as your feature and target Data
X_train_full, X_test, y_train_full, y_test = train_test_split(X, y,

# Define regression models
models = {
    'Linear Regression': LinearRegression(),
    'Decision Tree': DecisionTreeRegressor(),
    'Random Forest': RandomForestRegressor(n_estimators=100),
    'k-NN': KNeighborsRegressor(n_neighbors=10),
    'Gradient Boosting': GradientBoostingRegressor(),
    'XGBoost': xgb.XGBRegressor(objective='reg:squarederror'),
    'CatBoost': CatBoostRegressor(verbose=0),
    'LightGBM': lgb.LGBMRegressor()
}

# Simulate federated learning with 4 clients
parts_X = np.array_split(X_train_full, 4)
parts_y = np.array_split(y_train_full, 4)

results = []
kf = KFold(n_splits=5, shuffle=True, random_state=42)

# Collect predictions from each client
client_predictions = {
    name: {col: [] for col in y_train_full.columns} for name in models
}

# Federated-style training
for i in range(4):
    X_train, y_train = parts_X[i], parts_y[i]
```

```

for name, model in models.items():
    for col in y_train.columns:
        model.fit(X_train, y_train[col])
        predictions = model.predict(X_test)
        client_predictions[name][col].append(predictions)

# Average predictions and compute metrics
for name, targets in client_predictions.items():
    print(f"Results for {name}:")
    for target_col, predictions in targets.items():
        average_prediction = np.mean(predictions, axis=0)
        mse = mean_squared_error(y_test[target_col], average_prediction)
        mae = mean_absolute_error(y_test[target_col], average_prediction)
        rmse = np.sqrt(mse)
        mape = np.mean(np.abs((y_test[target_col] - average_prediction) / y_test[target_col]))
        r2 = r2_score(y_test[target_col], average_prediction)
        rse = np.sqrt(np.sum((y_test[target_col] - average_prediction) ** 2) / len(y_test[target_col]))

    print(f" {target_col} - Averaged MSE: {mse:.4f}, MAE: {mae:.4f}, RMSE: {rmse:.4f}, MAPE: {mape:.4f}, R2: {r2:.4f}, RSE: {rse:.4f}")

    results.append({
        'model': name,
        'target_col': target_col,
        'mse': mse,
        'mae': mae,
        'rmse': rmse,
        'mape': mape,
        'r2': r2,
        'rse': rse
    })

# Classical centralized training
for target_col in y_train_full.columns:
    print(f"\nResults for target variable: {target_col}")
    for name, model in models.items():
        start_time = time.time()
        model.fit(X_train_full, y_train_full[target_col])
        y_pred = model.predict(X_test)
        end_time = time.time()
        elapsed_time = end_time - start_time

        mse = mean_squared_error(y_test[target_col], y_pred)
        mae = mean_absolute_error(y_test[target_col], y_pred)
        rmse = np.sqrt(mse)
        mape = np.mean(np.abs((y_test[target_col] - y_pred) / y_test[target_col]))
        r2 = r2_score(y_test[target_col], y_pred)
        rse = np.sqrt(np.sum((y_test[target_col] - y_pred) ** 2) / len(y_test[target_col]))

    print(f"{name} - MSE: {mse:.4f}, MAE: {mae:.4f}, RMSE: {rmse:.4f}, MAPE: {mape:.4f}, R2: {r2:.4f}, RSE: {rse:.4f}, Elapsed Time: {elapsed_time:.4f}")

    results.append({
        'model': name,
        'target_col': target_col,
        'mse': mse,
        'mae': mae,
        'rmse': rmse,
        'mape': mape,
        'r2': r2,
        'rse': rse,
        'elapsed_time': elapsed_time
    })

```

```
        'mape': mape,
        'r2': r2,
        'rse': rse,
        'time': elapsed_time
    })

# Convert to DataFrame
results_df = pd.DataFrame(results)

# Plotting results
plt.figure(figsize=(14, 10))

# MSE
plt.subplot(2, 2, 1)
for name in results_df['model'].unique():
    model_results = results_df[results_df['model'] == name]
    plt.plot(model_results['target_col'], model_results['mse'], label=name)
plt.title('Mean Squared Error')
plt.xlabel('Target Variable')
plt.ylabel('MSE')
plt.xticks(rotation=45)
plt.legend()

# R^2
plt.subplot(2, 2, 2)
for name in results_df['model'].unique():
    model_results = results_df[results_df['model'] == name]
    plt.plot(model_results['target_col'], model_results['r2'], label=name)
plt.title('R^2 Score')
plt.xlabel('Target Variable')
plt.ylabel('R^2')
plt.xticks(rotation=45)
plt.legend()

# Time
plt.subplot(2, 2, 3)
for name in results_df['model'].unique():
    model_results = results_df[results_df['model'] == name]
    if 'time' in model_results.columns:
        plt.plot(model_results['target_col'], model_results['time'], label=name)
plt.title('Time Taken')
plt.xlabel('Target Variable')
plt.ylabel('Time (s)')
plt.xticks(rotation=45)
plt.legend()

# RMSE
plt.subplot(2, 2, 4)
for name in results_df['model'].unique():
    model_results = results_df[results_df['model'] == name]
    plt.plot(model_results['target_col'], model_results['rmse'], label=name)
plt.title('Root Mean Squared Error')
plt.xlabel('Target Variable')
plt.ylabel('RMSE')
plt.xticks(rotation=45)
plt.legend()
```



```
plt.tight_layout()  
plt.show()
```

```
/usr/local/lib/python3.11/dist-packages/numpy/_core/fromnumeric.py:5  
7: FutureWarning: 'DataFrame.swapaxes' is deprecated and will be rem  
oved in a future version. Please use 'DataFrame.transpose' instead.
```

```
    return bound(*args, **kwargs)
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhea  
d of testing was 0.011217 seconds.
```

```
You can set `force_row_wise=true` to remove the overhead.
```

```
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 1099
```

```
[LightGBM] [Info] Number of data points in the train set: 107573, nu  
mber of used features: 13
```

```
[LightGBM] [Info] Start training from score 2.560634
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhea  
d of testing was 0.017887 seconds.
```

```
You can set `force_row_wise=true` to remove the overhead.
```

```
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 1099
```

```
[LightGBM] [Info] Number of data points in the train set: 107573, nu  
mber of used features: 13
```

```
[LightGBM] [Info] Start training from score 2.582565
```

```
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhea  
d of testing was 0.033130 seconds.
```

```
You can set `force_col_wise=true` to remove the overhead.
```

```
[LightGBM] [Info] Total Bins 1099
```

```
[LightGBM] [Info] Number of data points in the train set: 107573, nu  
mber of used features: 13
```

```
[LightGBM] [Info] Start training from score 2.652507
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhea  
d of testing was 0.010569 seconds.
```

```
You can set `force_row_wise=true` to remove the overhead.
```

```
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 1099
```

```
[LightGBM] [Info] Number of data points in the train set: 107573, nu  
mber of used features: 13
```

```
[LightGBM] [Info] Start training from score 2.680321
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhea  
d of testing was 0.010523 seconds.
```

```
You can set `force_row_wise=true` to remove the overhead.
```

```
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 1099
```

```
[LightGBM] [Info] Number of data points in the train set: 107573, nu  
mber of used features: 13
```

```
[LightGBM] [Info] Start training from score 2.637388
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhea  
d of testing was 0.010375 seconds.
```

```
You can set `force_row_wise=true` to remove the overhead.
```

```
And if memory is not enough, you can set `force_col_wise=true`.
```

```
[LightGBM] [Info] Total Bins 1099
```

```
[LightGBM] [Info] Number of data points in the train set: 107573, nu  
mber of used features: 13
```

```
[LightGBM] [Info] Start training from score 2.664167
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhea
```

ad of testing was 0.012354 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 3.347749
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010640 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 3.431678
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010609 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.269855
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.015694 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.302631
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008610 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.560517
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.015057 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.582592
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.009560 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.653578
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead

ad of testing was 0.008766 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.680966
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008750 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.638017
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008430 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.663916
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008622 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 3.347473
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008694 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 3.432815
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.013348 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.271818
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008766 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.303505
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead

ad of testing was 0.011562 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.559775
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010471 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.581604
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010682 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.652370
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.019319 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.679292
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010532 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.637137
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.015479 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.663156
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010523 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 3.346177
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010567 seconds.

You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 3.430175
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010562 seconds.

You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.269549
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.010490 seconds.

You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1101
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.302539
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008772 seconds.

You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.560990
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008705 seconds.

You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.583033
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.014160 seconds.

You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.652748
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008734 seconds.

You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.680719
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008882 seconds.

You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.637060
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.019909 seconds.
You can set `force_col_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.663635
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.009162 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 3.347197
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008703 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 3.431502
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.008807 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.271324
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.013187 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1099
[LightGBM] [Info] Number of data points in the train set: 107573, number of used features: 13
[LightGBM] [Info] Start training from score 2.303219

Results for Linear Regression:

ST58TA - Averaged MSE: 0.0287, MAE: 0.1311, RMSE: 0.1695, MAPE: 5.27%, R2: 0.2937, RSE: 0.1695
ST117TA - Averaged MSE: 0.0287, MAE: 0.1303, RMSE: 0.1695, MAPE: 5.15%, R2: 0.3011, RSE: 0.1695
ST40TA - Averaged MSE: 0.0287, MAE: 0.1273, RMSE: 0.1693, MAPE: 4.95%, R2: 0.1904, RSE: 0.1693
ST99TA - Averaged MSE: 0.0277, MAE: 0.1256, RMSE: 0.1664, MAPE: 4.81%, R2: 0.1649, RSE: 0.1664
ST32TA - Averaged MSE: 0.0302, MAE: 0.1328, RMSE: 0.1737, MAPE: 5.16%, R2: 0.1293, RSE: 0.1737

ST91TA – Averaged MSE: 0.0282, MAE: 0.1288, RMSE: 0.1680, MAPE: 4.98%, R2: 0.1500, RSE: 0.1680

ST60TA – Averaged MSE: 0.1342, MAE: 0.2705, RMSE: 0.3664, MAPE: 8.82%, R2: 0.0553, RSE: 0.3664

ST119TA – Averaged MSE: 0.1285, MAE: 0.2619, RMSE: 0.3585, MAPE: 8.41%, R2: 0.1191, RSE: 0.3585

ST62TA – Averaged MSE: 0.0389, MAE: 0.1573, RMSE: 0.1974, MAPE: 7.10%, R2: 0.1289, RSE: 0.1974

ST121TA – Averaged MSE: 0.0398, MAE: 0.1570, RMSE: 0.1995, MAPE: 6.94%, R2: 0.1548, RSE: 0.1995

Results for Decision Tree:

ST58TA – Averaged MSE: 0.0042, MAE: 0.0484, RMSE: 0.0651, MAPE: 1.91%, R2: 0.8960, RSE: 0.0651

ST117TA – Averaged MSE: 0.0038, MAE: 0.0444, RMSE: 0.0613, MAPE: 1.74%, R2: 0.9087, RSE: 0.0613

ST40TA – Averaged MSE: 0.0043, MAE: 0.0468, RMSE: 0.0654, MAPE: 1.78%, R2: 0.8792, RSE: 0.0654

ST99TA – Averaged MSE: 0.0039, MAE: 0.0447, RMSE: 0.0625, MAPE: 1.70%, R2: 0.8820, RSE: 0.0625

ST32TA – Averaged MSE: 0.0053, MAE: 0.0527, RMSE: 0.0728, MAPE: 2.01%, R2: 0.8473, RSE: 0.0728

ST91TA – Averaged MSE: 0.0059, MAE: 0.0549, RMSE: 0.0769, MAPE: 2.11%, R2: 0.8217, RSE: 0.0769

ST60TA – Averaged MSE: 0.0276, MAE: 0.1221, RMSE: 0.1663, MAPE: 3.83%, R2: 0.8054, RSE: 0.1663

ST119TA – Averaged MSE: 0.0324, MAE: 0.1314, RMSE: 0.1801, MAPE: 3.96%, R2: 0.7778, RSE: 0.1801

ST62TA – Averaged MSE: 0.0067, MAE: 0.0615, RMSE: 0.0820, MAPE: 2.74%, R2: 0.8497, RSE: 0.0820

ST121TA – Averaged MSE: 0.0081, MAE: 0.0664, RMSE: 0.0898, MAPE: 2.92%, R2: 0.8287, RSE: 0.0898

Results for Random Forest:

ST58TA – Averaged MSE: 0.0042, MAE: 0.0482, RMSE: 0.0648, MAPE: 1.91%, R2: 0.8969, RSE: 0.0648

ST117TA – Averaged MSE: 0.0037, MAE: 0.0442, RMSE: 0.0610, MAPE: 1.73%, R2: 0.9096, RSE: 0.0610

ST40TA – Averaged MSE: 0.0042, MAE: 0.0466, RMSE: 0.0650, MAPE: 1.78%, R2: 0.8805, RSE: 0.0650

ST99TA – Averaged MSE: 0.0039, MAE: 0.0446, RMSE: 0.0622, MAPE: 1.69%, R2: 0.8831, RSE: 0.0622

ST32TA – Averaged MSE: 0.0052, MAE: 0.0525, RMSE: 0.0724, MAPE: 2.00%, R2: 0.8490, RSE: 0.0724

ST91TA – Averaged MSE: 0.0059, MAE: 0.0547, RMSE: 0.0765, MAPE: 2.10%, R2: 0.8235, RSE: 0.0765

ST60TA – Averaged MSE: 0.0273, MAE: 0.1216, RMSE: 0.1653, MAPE: 3.81%, R2: 0.8078, RSE: 0.1653

ST119TA – Averaged MSE: 0.0321, MAE: 0.1310, RMSE: 0.1792, MAPE: 3.95%, R2: 0.7800, RSE: 0.1792

ST62TA – Averaged MSE: 0.0067, MAE: 0.0612, RMSE: 0.0816, MAPE: 2.73%, R2: 0.8511, RSE: 0.0816

ST121TA – Averaged MSE: 0.0080, MAE: 0.0661, RMSE: 0.0894, MAPE: 2.91%, R2: 0.8304, RSE: 0.0894

Results for k-NN:

ST58TA – Averaged MSE: 0.0045, MAE: 0.0494, RMSE: 0.0674, MAPE: 1.96%, R2: 0.8884, RSE: 0.0674

ST117TA – Averaged MSE: 0.0040, MAE: 0.0453, RMSE: 0.0634, MAPE: 1.

78%, R2: 0.9021, RSE: 0.0634
ST40TA – Averaged MSE: 0.0045, MAE: 0.0475, RMSE: 0.0672, MAPE: 1.8
1%, R2: 0.8724, RSE: 0.0672
ST99TA – Averaged MSE: 0.0041, MAE: 0.0455, RMSE: 0.0642, MAPE: 1.7
3%, R2: 0.8758, RSE: 0.0642
ST32TA – Averaged MSE: 0.0055, MAE: 0.0533, RMSE: 0.0742, MAPE: 2.0
3%, R2: 0.8413, RSE: 0.0742
ST91TA – Averaged MSE: 0.0061, MAE: 0.0555, RMSE: 0.0781, MAPE: 2.1
3%, R2: 0.8162, RSE: 0.0781
ST60TA – Averaged MSE: 0.0284, MAE: 0.1233, RMSE: 0.1687, MAPE: 3.8
7%, R2: 0.7998, RSE: 0.1687
ST119TA – Averaged MSE: 0.0335, MAE: 0.1326, RMSE: 0.1829, MAPE: 4.
02%, R2: 0.7707, RSE: 0.1829
ST62TA – Averaged MSE: 0.0070, MAE: 0.0624, RMSE: 0.0839, MAPE: 2.7
9%, R2: 0.8427, RSE: 0.0839
ST121TA – Averaged MSE: 0.0084, MAE: 0.0673, RMSE: 0.0917, MAPE: 2.
97%, R2: 0.8216, RSE: 0.0917
Results for Gradient Boosting:
ST58TA – Averaged MSE: 0.0141, MAE: 0.0923, RMSE: 0.1188, MAPE: 3.7
0%, R2: 0.6533, RSE: 0.1188
ST117TA – Averaged MSE: 0.0128, MAE: 0.0869, RMSE: 0.1132, MAPE: 3.
45%, R2: 0.6879, RSE: 0.1132
ST40TA – Averaged MSE: 0.0143, MAE: 0.0898, RMSE: 0.1197, MAPE: 3.4
8%, R2: 0.5954, RSE: 0.1197
ST99TA – Averaged MSE: 0.0131, MAE: 0.0855, RMSE: 0.1144, MAPE: 3.2
8%, R2: 0.6053, RSE: 0.1144
ST32TA – Averaged MSE: 0.0161, MAE: 0.0958, RMSE: 0.1269, MAPE: 3.7
0%, R2: 0.5355, RSE: 0.1269
ST91TA – Averaged MSE: 0.0153, MAE: 0.0914, RMSE: 0.1236, MAPE: 3.5
3%, R2: 0.5394, RSE: 0.1236
ST60TA – Averaged MSE: 0.0809, MAE: 0.2106, RMSE: 0.2845, MAPE: 6.7
9%, R2: 0.4304, RSE: 0.2845
ST119TA – Averaged MSE: 0.0848, MAE: 0.2146, RMSE: 0.2911, MAPE: 6.
81%, R2: 0.4190, RSE: 0.2911
ST62TA – Averaged MSE: 0.0219, MAE: 0.1151, RMSE: 0.1478, MAPE: 5.1
8%, R2: 0.5111, RSE: 0.1478
ST121TA – Averaged MSE: 0.0239, MAE: 0.1170, RMSE: 0.1544, MAPE: 5.
18%, R2: 0.4934, RSE: 0.1544
Results for XGBoost:
ST58TA – Averaged MSE: 0.0043, MAE: 0.0493, RMSE: 0.0656, MAPE: 1.9
5%, R2: 0.8942, RSE: 0.0656
ST117TA – Averaged MSE: 0.0039, MAE: 0.0456, RMSE: 0.0621, MAPE: 1.
79%, R2: 0.9061, RSE: 0.0621
ST40TA – Averaged MSE: 0.0043, MAE: 0.0478, RMSE: 0.0659, MAPE: 1.8
2%, R2: 0.8773, RSE: 0.0659
ST99TA – Averaged MSE: 0.0040, MAE: 0.0461, RMSE: 0.0634, MAPE: 1.7
5%, R2: 0.8787, RSE: 0.0634
ST32TA – Averaged MSE: 0.0054, MAE: 0.0537, RMSE: 0.0732, MAPE: 2.0
5%, R2: 0.8457, RSE: 0.0732
ST91TA – Averaged MSE: 0.0059, MAE: 0.0556, RMSE: 0.0768, MAPE: 2.1
3%, R2: 0.8224, RSE: 0.0768
ST60TA – Averaged MSE: 0.0278, MAE: 0.1239, RMSE: 0.1668, MAPE: 3.8
9%, R2: 0.8041, RSE: 0.1668
ST119TA – Averaged MSE: 0.0326, MAE: 0.1334, RMSE: 0.1805, MAPE: 4.
04%, R2: 0.7767, RSE: 0.1805
ST62TA – Averaged MSE: 0.0069, MAE: 0.0627, RMSE: 0.0828, MAPE: 2.8

0%, R2: 0.8466, RSE: 0.0828

ST121TA – Averaged MSE: 0.0083, MAE: 0.0679, RMSE: 0.0910, MAPE: 2.99%, R2: 0.8240, RSE: 0.0910

Results for CatBoost:

ST58TA – Averaged MSE: 0.0045, MAE: 0.0506, RMSE: 0.0673, MAPE: 2.01%, R2: 0.8886, RSE: 0.0673

ST117TA – Averaged MSE: 0.0040, MAE: 0.0466, RMSE: 0.0632, MAPE: 1.83%, R2: 0.9027, RSE: 0.0632

ST40TA – Averaged MSE: 0.0045, MAE: 0.0486, RMSE: 0.0669, MAPE: 1.86%, R2: 0.8736, RSE: 0.0669

ST99TA – Averaged MSE: 0.0041, MAE: 0.0468, RMSE: 0.0642, MAPE: 1.78%, R2: 0.8757, RSE: 0.0642

ST32TA – Averaged MSE: 0.0055, MAE: 0.0547, RMSE: 0.0741, MAPE: 2.09%, R2: 0.8415, RSE: 0.0741

ST91TA – Averaged MSE: 0.0060, MAE: 0.0567, RMSE: 0.0777, MAPE: 2.18%, R2: 0.8179, RSE: 0.0777

ST60TA – Averaged MSE: 0.0284, MAE: 0.1253, RMSE: 0.1684, MAPE: 3.93%, R2: 0.8004, RSE: 0.1684

ST119TA – Averaged MSE: 0.0332, MAE: 0.1351, RMSE: 0.1821, MAPE: 4.10%, R2: 0.7728, RSE: 0.1821

ST62TA – Averaged MSE: 0.0071, MAE: 0.0640, RMSE: 0.0844, MAPE: 2.86%, R2: 0.8408, RSE: 0.0844

ST121TA – Averaged MSE: 0.0086, MAE: 0.0691, RMSE: 0.0926, MAPE: 3.05%, R2: 0.8180, RSE: 0.0926

Results for LightGBM:

ST58TA – Averaged MSE: 0.0060, MAE: 0.0586, RMSE: 0.0777, MAPE: 2.33%, R2: 0.8517, RSE: 0.0777

ST117TA – Averaged MSE: 0.0055, MAE: 0.0546, RMSE: 0.0739, MAPE: 2.15%, R2: 0.8671, RSE: 0.0739

ST40TA – Averaged MSE: 0.0060, MAE: 0.0565, RMSE: 0.0772, MAPE: 2.16%, R2: 0.8318, RSE: 0.0772

ST99TA – Averaged MSE: 0.0056, MAE: 0.0546, RMSE: 0.0745, MAPE: 2.08%, R2: 0.8323, RSE: 0.0745

ST32TA – Averaged MSE: 0.0071, MAE: 0.0625, RMSE: 0.0841, MAPE: 2.39%, R2: 0.7958, RSE: 0.0841

ST91TA – Averaged MSE: 0.0077, MAE: 0.0645, RMSE: 0.0877, MAPE: 2.48%, R2: 0.7684, RSE: 0.0877

ST60TA – Averaged MSE: 0.0354, MAE: 0.1405, RMSE: 0.1881, MAPE: 4.42%, R2: 0.7511, RSE: 0.1881

ST119TA – Averaged MSE: 0.0405, MAE: 0.1502, RMSE: 0.2013, MAPE: 4.60%, R2: 0.7222, RSE: 0.2013

ST62TA – Averaged MSE: 0.0097, MAE: 0.0756, RMSE: 0.0987, MAPE: 3.38%, R2: 0.7820, RSE: 0.0987

ST121TA – Averaged MSE: 0.0117, MAE: 0.0808, RMSE: 0.1082, MAPE: 3.58%, R2: 0.7512, RSE: 0.1082

Results for target variable: ST58TA

Linear Regression – MSE: 0.0287, MAE: 0.1311, RMSE: 0.1695, MAPE: 5.27%, R2: 0.2937, RSE: 0.1695, Time: 0.22s

Decision Tree – MSE: 0.0042, MAE: 0.0484, RMSE: 0.0651, MAPE: 1.92%, R2: 0.8959, RSE: 0.0651, Time: 1.67s

Random Forest – MSE: 0.0042, MAE: 0.0484, RMSE: 0.0650, MAPE: 1.91%, R2: 0.8960, RSE: 0.0650, Time: 127.58s

k-NN – MSE: 0.0045, MAE: 0.0494, RMSE: 0.0672, MAPE: 1.95%, R2: 0.8891, RSE: 0.0672, Time: 7.54s

Gradient Boosting – MSE: 0.0138, MAE: 0.0913, RMSE: 0.1174, MAPE: 3.

66%, R2: 0.6614, RSE: 0.1174, Time: 56.03s
XGBoost – MSE: 0.0043, MAE: 0.0495, RMSE: 0.0659, MAPE: 1.96%, R2: 0.8933, RSE: 0.0659, Time: 4.97s
CatBoost – MSE: 0.0043, MAE: 0.0494, RMSE: 0.0657, MAPE: 1.96%, R2: 0.8939, RSE: 0.0657, Time: 62.29s
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.061071 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.560479
LightGBM – MSE: 0.0061, MAE: 0.0588, RMSE: 0.0781, MAPE: 2.33%, R2: 0.8501, RSE: 0.0781, Time: 5.03s

Results for target variable: ST117TA

Linear Regression – MSE: 0.0287, MAE: 0.1303, RMSE: 0.1695, MAPE: 5.15%, R2: 0.3011, RSE: 0.1695, Time: 0.19s
Decision Tree – MSE: 0.0037, MAE: 0.0444, RMSE: 0.0612, MAPE: 1.74%, R2: 0.9090, RSE: 0.0612, Time: 1.58s
Random Forest – MSE: 0.0037, MAE: 0.0443, RMSE: 0.0611, MAPE: 1.74%, R2: 0.9090, RSE: 0.0611, Time: 122.47s
k-NN – MSE: 0.0040, MAE: 0.0455, RMSE: 0.0631, MAPE: 1.78%, R2: 0.9031, RSE: 0.0631, Time: 7.45s
Gradient Boosting – MSE: 0.0129, MAE: 0.0872, RMSE: 0.1137, MAPE: 3.47%, R2: 0.6856, RSE: 0.1137, Time: 56.76s
XGBoost – MSE: 0.0039, MAE: 0.0457, RMSE: 0.0622, MAPE: 1.80%, R2: 0.9057, RSE: 0.0622, Time: 3.19s
CatBoost – MSE: 0.0038, MAE: 0.0457, RMSE: 0.0619, MAPE: 1.79%, R2: 0.9067, RSE: 0.0619, Time: 65.27s
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.061330 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.582449
LightGBM – MSE: 0.0054, MAE: 0.0547, RMSE: 0.0738, MAPE: 2.16%, R2: 0.8674, RSE: 0.0738, Time: 5.59s

Results for target variable: ST40TA

Linear Regression – MSE: 0.0287, MAE: 0.1273, RMSE: 0.1693, MAPE: 4.95%, R2: 0.1904, RSE: 0.1693, Time: 0.20s
Decision Tree – MSE: 0.0043, MAE: 0.0468, RMSE: 0.0654, MAPE: 1.78%, R2: 0.8793, RSE: 0.0654, Time: 1.65s
Random Forest – MSE: 0.0043, MAE: 0.0468, RMSE: 0.0653, MAPE: 1.78%, R2: 0.8794, RSE: 0.0653, Time: 124.90s
k-NN – MSE: 0.0045, MAE: 0.0479, RMSE: 0.0674, MAPE: 1.83%, R2: 0.8716, RSE: 0.0674, Time: 7.80s
Gradient Boosting – MSE: 0.0145, MAE: 0.0901, RMSE: 0.1203, MAPE: 3.49%, R2: 0.5910, RSE: 0.1203, Time: 56.96s
XGBoost – MSE: 0.0044, MAE: 0.0480, RMSE: 0.0662, MAPE: 1.83%, R2: 0.8761, RSE: 0.0662, Time: 5.03s
CatBoost – MSE: 0.0043, MAE: 0.0476, RMSE: 0.0655, MAPE: 1.82%, R2:

0.8788, RSE: 0.0655, Time: 63.10s
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.049409 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.652801
LightGBM - MSE: 0.0060, MAE: 0.0567, RMSE: 0.0773, MAPE: 2.17%, R2: 0.8313, RSE: 0.0773, Time: 4.82s

Results for target variable: ST99TA

Linear Regression - MSE: 0.0277, MAE: 0.1256, RMSE: 0.1664, MAPE: 4.81%, R2: 0.1649, RSE: 0.1664, Time: 0.21s
Decision Tree - MSE: 0.0039, MAE: 0.0447, RMSE: 0.0625, MAPE: 1.70%, R2: 0.8823, RSE: 0.0625, Time: 1.69s
Random Forest - MSE: 0.0039, MAE: 0.0447, RMSE: 0.0624, MAPE: 1.70%, R2: 0.8824, RSE: 0.0624, Time: 120.69s
k-NN - MSE: 0.0042, MAE: 0.0457, RMSE: 0.0645, MAPE: 1.74%, R2: 0.8744, RSE: 0.0645, Time: 7.17s
Gradient Boosting - MSE: 0.0131, MAE: 0.0857, RMSE: 0.1147, MAPE: 3.29%, R2: 0.6033, RSE: 0.1147, Time: 56.30s
XGBoost - MSE: 0.0041, MAE: 0.0465, RMSE: 0.0637, MAPE: 1.77%, R2: 0.8776, RSE: 0.0637, Time: 3.14s
CatBoost - MSE: 0.0040, MAE: 0.0458, RMSE: 0.0629, MAPE: 1.74%, R2: 0.8808, RSE: 0.0629, Time: 63.20s
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.044250 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.680324
LightGBM - MSE: 0.0056, MAE: 0.0552, RMSE: 0.0751, MAPE: 2.10%, R2: 0.8296, RSE: 0.0751, Time: 5.09s

Results for target variable: ST32TA

Linear Regression - MSE: 0.0302, MAE: 0.1328, RMSE: 0.1737, MAPE: 5.16%, R2: 0.1293, RSE: 0.1737, Time: 0.40s
Decision Tree - MSE: 0.0053, MAE: 0.0526, RMSE: 0.0727, MAPE: 2.01%, R2: 0.8476, RSE: 0.0727, Time: 1.91s
Random Forest - MSE: 0.0053, MAE: 0.0526, RMSE: 0.0726, MAPE: 2.01%, R2: 0.8478, RSE: 0.0726, Time: 123.87s
k-NN - MSE: 0.0056, MAE: 0.0536, RMSE: 0.0748, MAPE: 2.04%, R2: 0.8386, RSE: 0.0748, Time: 7.11s
Gradient Boosting - MSE: 0.0159, MAE: 0.0950, RMSE: 0.1261, MAPE: 3.67%, R2: 0.5414, RSE: 0.1261, Time: 59.32s
XGBoost - MSE: 0.0054, MAE: 0.0539, RMSE: 0.0733, MAPE: 2.06%, R2: 0.8450, RSE: 0.0733, Time: 3.24s
CatBoost - MSE: 0.0053, MAE: 0.0535, RMSE: 0.0728, MAPE: 2.04%, R2: 0.8473, RSE: 0.0728, Time: 63.53s
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.042766 seconds.
You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.637400
LightGBM – MSE: 0.0071, MAE: 0.0630, RMSE: 0.0845, MAPE: 2.41%, R2: 0.7942, RSE: 0.0845, Time: 5.58s

Results for target variable: ST91TA
Linear Regression – MSE: 0.0282, MAE: 0.1288, RMSE: 0.1680, MAPE: 4.98%, R2: 0.1500, RSE: 0.1680, Time: 0.22s
Decision Tree – MSE: 0.0059, MAE: 0.0549, RMSE: 0.0768, MAPE: 2.11%, R2: 0.8221, RSE: 0.0768, Time: 1.61s
Random Forest – MSE: 0.0059, MAE: 0.0548, RMSE: 0.0768, MAPE: 2.10%, R2: 0.8222, RSE: 0.0768, Time: 120.79s
k-NN – MSE: 0.0063, MAE: 0.0562, RMSE: 0.0793, MAPE: 2.15%, R2: 0.8104, RSE: 0.0793, Time: 7.34s
Gradient Boosting – MSE: 0.0153, MAE: 0.0916, RMSE: 0.1239, MAPE: 3.53%, R2: 0.5377, RSE: 0.1239, Time: 59.07s
XGBoost – MSE: 0.0060, MAE: 0.0561, RMSE: 0.0772, MAPE: 2.15%, R2: 0.8204, RSE: 0.0772, Time: 3.10s
CatBoost – MSE: 0.0058, MAE: 0.0556, RMSE: 0.0765, MAPE: 2.13%, R2: 0.8238, RSE: 0.0765, Time: 65.07s
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.042317 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.663719
LightGBM – MSE: 0.0078, MAE: 0.0651, RMSE: 0.0884, MAPE: 2.50%, R2: 0.7648, RSE: 0.0884, Time: 5.03s

Results for target variable: ST60TA
Linear Regression – MSE: 0.1342, MAE: 0.2705, RMSE: 0.3664, MAPE: 8.82%, R2: 0.0553, RSE: 0.3664, Time: 0.34s
Decision Tree – MSE: 0.0277, MAE: 0.1221, RMSE: 0.1663, MAPE: 3.83%, R2: 0.8054, RSE: 0.1663, Time: 2.18s
Random Forest – MSE: 0.0276, MAE: 0.1220, RMSE: 0.1661, MAPE: 3.82%, R2: 0.8057, RSE: 0.1661, Time: 146.28s
k-NN – MSE: 0.0293, MAE: 0.1250, RMSE: 0.1710, MAPE: 3.91%, R2: 0.7941, RSE: 0.1710, Time: 8.37s
Gradient Boosting – MSE: 0.0816, MAE: 0.2111, RMSE: 0.2856, MAPE: 6.82%, R2: 0.4259, RSE: 0.2856, Time: 58.03s
XGBoost – MSE: 0.0282, MAE: 0.1250, RMSE: 0.1678, MAPE: 3.92%, R2: 0.8018, RSE: 0.1678, Time: 3.15s
CatBoost – MSE: 0.0274, MAE: 0.1233, RMSE: 0.1656, MAPE: 3.87%, R2: 0.8069, RSE: 0.1656, Time: 63.54s
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.042176 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 3.347149
LightGBM – MSE: 0.0358, MAE: 0.1413, RMSE: 0.1892, MAPE: 4.45%, R2:
0.7481, RSE: 0.1892, Time: 4.79s

Results for target variable: ST119TA

Linear Regression – MSE: 0.1285, MAE: 0.2619, RMSE: 0.3585, MAPE: 8.41%, R2: 0.1191, RSE: 0.3585, Time: 0.34s

Decision Tree – MSE: 0.0325, MAE: 0.1315, RMSE: 0.1803, MAPE: 3.97%, R2: 0.7773, RSE: 0.1803, Time: 2.75s

Random Forest – MSE: 0.0324, MAE: 0.1314, RMSE: 0.1801, MAPE: 3.96%, R2: 0.7776, RSE: 0.1801, Time: 137.84s

k-NN – MSE: 0.0343, MAE: 0.1341, RMSE: 0.1853, MAPE: 4.04%, R2: 0.7648, RSE: 0.1853, Time: 6.88s

Gradient Boosting – MSE: 0.0846, MAE: 0.2145, RMSE: 0.2909, MAPE: 6.80%, R2: 0.4198, RSE: 0.2909, Time: 59.03s

XGBoost – MSE: 0.0326, MAE: 0.1335, RMSE: 0.1804, MAPE: 4.04%, R2: 0.7769, RSE: 0.1804, Time: 3.19s

CatBoost – MSE: 0.0322, MAE: 0.1327, RMSE: 0.1793, MAPE: 4.02%, R2: 0.7795, RSE: 0.1793, Time: 63.42s

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.077088 seconds.

You can set `force_col_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 3.431542

LightGBM – MSE: 0.0405, MAE: 0.1503, RMSE: 0.2013, MAPE: 4.60%, R2: 0.7222, RSE: 0.2013, Time: 6.67s

Results for target variable: ST62TA

Linear Regression – MSE: 0.0389, MAE: 0.1573, RMSE: 0.1974, MAPE: 7.10%, R2: 0.1289, RSE: 0.1974, Time: 0.20s

Decision Tree – MSE: 0.0067, MAE: 0.0615, RMSE: 0.0820, MAPE: 2.74%, R2: 0.8496, RSE: 0.0820, Time: 1.75s

Random Forest – MSE: 0.0067, MAE: 0.0615, RMSE: 0.0819, MAPE: 2.74%, R2: 0.8499, RSE: 0.0819, Time: 130.48s

k-NN – MSE: 0.0071, MAE: 0.0629, RMSE: 0.0844, MAPE: 2.81%, R2: 0.8406, RSE: 0.0844, Time: 7.74s

Gradient Boosting – MSE: 0.0218, MAE: 0.1147, RMSE: 0.1476, MAPE: 5.17%, R2: 0.5126, RSE: 0.1476, Time: 57.35s

XGBoost – MSE: 0.0069, MAE: 0.0630, RMSE: 0.0831, MAPE: 2.81%, R2: 0.8456, RSE: 0.0831, Time: 3.70s

CatBoost – MSE: 0.0068, MAE: 0.0624, RMSE: 0.0824, MAPE: 2.79%, R2: 0.8483, RSE: 0.0824, Time: 65.09s

[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.042551 seconds.

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

[LightGBM] [Info] Total Bins 1102


[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.270637

LightGBM – MSE: 0.0099, MAE: 0.0762, RMSE: 0.0995, MAPE: 3.41%, R2: 0.7786, RSE: 0.0995, Time: 4.58s

Results for target variable: ST121TA

Linear Regression – MSE: 0.0398, MAE: 0.1570, RMSE: 0.1995, MAPE: 6.94%, R2: 0.1548, RSE: 0.1995, Time: 0.19s
 Decision Tree – MSE: 0.0081, MAE: 0.0665, RMSE: 0.0899, MAPE: 2.93%, R2: 0.8283, RSE: 0.0899, Time: 1.73s
 Random Forest – MSE: 0.0081, MAE: 0.0664, RMSE: 0.0898, MAPE: 2.92%, R2: 0.8286, RSE: 0.0898, Time: 130.98s
 k-NN – MSE: 0.0086, MAE: 0.0681, RMSE: 0.0926, MAPE: 3.00%, R2: 0.8179, RSE: 0.0926, Time: 7.15s
 Gradient Boosting – MSE: 0.0238, MAE: 0.1168, RMSE: 0.1543, MAPE: 5.17%, R2: 0.4941, RSE: 0.1543, Time: 59.76s
 XGBoost – MSE: 0.0082, MAE: 0.0677, RMSE: 0.0906, MAPE: 2.98%, R2: 0.8256, RSE: 0.0906, Time: 5.07s
 CatBoost – MSE: 0.0082, MAE: 0.0677, RMSE: 0.0906, MAPE: 2.98%, R2: 0.8259, RSE: 0.0906, Time: 63.94s
 [LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.043421 seconds.
 You can set `force_row_wise=true` to remove the overhead.
 And if memory is not enough, you can set `force_col_wise=true`.
 [LightGBM] [Info] Total Bins 1102
 [LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
 [LightGBM] [Info] Start training from score 2.302973
 LightGBM – MSE: 0.0121, MAE: 0.0824, RMSE: 0.1099, MAPE: 3.64%, R2: 0.7435, RSE: 0.1099, Time: 4.52s



```
In [ ]: import pandas as pd
import numpy as np
import time
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.neighbors import KNeighborsRegressor
```

```

from sklearn.metrics import mean_squared_error, r2_score, mean_absolu
import xgboost as xgb
from catboost import CatBoostRegressor
import lightgbm as lgb
import matplotlib.pyplot as plt

# Assuming 'X' and 'y' are defined before this point
# Example:
# X = df.drop(columns=target_columns)
# y = df[target_columns] # where target_columns is a list of target

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=

# Initialize base models
base_models = {
    'Linear Regression': LinearRegression(),
    'Decision Tree': DecisionTreeRegressor(),
    'Random Forest': RandomForestRegressor(n_estimators=100),
    'k-NN': KNeighborsRegressor(n_neighbors=10),
    'Gradient Boosting': GradientBoostingRegressor(),
    'XGBoost': xgb.XGBRegressor(objective='reg:squarederror'),
    'CatBoost': CatBoostRegressor(verbose=0),
    'LightGBM': lgb.LGBMRegressor(force_col_wise=True)
}

results = []

# Train and evaluate each base model
for target_col in y_train.columns:
    print(f"Results for target variable: {target_col}")
    for name, model in base_models.items():
        start_time = time.time()
        model.fit(X_train, y_train[target_col])
        y_pred = model.predict(X_test)
        end_time = time.time()

        mse = mean_squared_error(y_test[target_col], y_pred)
        mae = mean_absolute_error(y_test[target_col], y_pred)
        rmse = np.sqrt(mse)
        mape = np.mean(np.abs((y_test[target_col] - y_pred) / y_test[target_col]))
        r2 = r2_score(y_test[target_col], y_pred)
        rse = np.sqrt(np.sum((y_test[target_col] - y_pred) ** 2) / len(y_test[target_col]))

        print(f"{name} - MSE: {mse:.4f}, MAE: {mae:.4f}, RMSE: {rmse:.4f}, R2: {r2:.4f}, RSE: {rse:.4f}")

    results.append({
        'model': name,
        'target_col': target_col,
        'mse': mse,
        'mae': mae,
        'rmse': rmse,
        'mape': mape,
        'r2': r2,
        'rse': rse,
        'time': end_time - start_time
    })

```

```

    })

# Create Voting Regressor
voting_regressor = VotingRegressor(estimators=[
    ('lr', base_models['Linear Regression']),
    ('dt', base_models['Decision Tree']),
    ('rf', base_models['Random Forest']),
    ('knn', base_models['k-NN']),
    ('gb', base_models['Gradient Boosting']),
    ('xgb', base_models['XGBoost']),
    ('catboost', base_models['CatBoost']),
    ('lgbm', base_models['LightGBM'])
])

# Train and evaluate Voting Regressor
for target_col in y_train.columns:
    print(f"\nVoting Regressor Results for target variable: {target_col}")
    start_time = time.time()
    voting_regressor.fit(X_train, y_train[target_col])
    y_pred = voting_regressor.predict(X_test)
    end_time = time.time()

    mse = mean_squared_error(y_test[target_col], y_pred)
    mae = mean_absolute_error(y_test[target_col], y_pred)
    rmse = np.sqrt(mse)
    mape = np.mean(np.abs((y_test[target_col] - y_pred) / y_test[target_col]))
    r2 = r2_score(y_test[target_col], y_pred)
    rse = np.sqrt(np.sum((y_test[target_col] - y_pred) ** 2) / (len(y_test) - 1))

    print(f"Voting Regressor - MSE: {mse:.4f}, MAE: {mae:.4f}, RMSE: {rmse:.4f}, R2: {r2:.4f}, RSE: {rse:.4f}")

    results.append({
        'model': 'Voting Regressor',
        'target_col': target_col,
        'mse': mse,
        'mae': mae,
        'rmse': rmse,
        'mape': mape,
        'r2': r2,
        'rse': rse,
        'time': end_time - start_time
    })

# Convert results to DataFrame
results_df = pd.DataFrame(results)

# Plotting results
for metric in ['mse', 'r2', 'time', 'rmse']:
    plt.figure(figsize=(12, 6))
    for name in results_df['model'].unique():
        model_results = results_df[results_df['model'] == name]
        plt.plot(model_results['target_col'], model_results[metric])
    plt.title(f'{metric.upper()} for each model across target variables')
    plt.xlabel('Target Variable')
    plt.ylabel(metric.upper())
    plt.legend()

```



```
plt.grid(True)
plt.tight_layout()
plt.show()
```

Results for target variable: ST58TA

Linear Regression – MSE: 0.0287, MAE: 0.1311, RMSE: 0.1695, MAPE: 5.27%, R2: 0.2937, Time: 0.23s

Decision Tree – MSE: 0.0042, MAE: 0.0484, RMSE: 0.0651, MAPE: 1.92%, R2: 0.8959, Time: 1.65s

Random Forest – MSE: 0.0042, MAE: 0.0484, RMSE: 0.0650, MAPE: 1.91%, R2: 0.8960, Time: 126.40s

k-NN – MSE: 0.0045, MAE: 0.0494, RMSE: 0.0672, MAPE: 1.95%, R2: 0.8891, Time: 7.72s

Gradient Boosting – MSE: 0.0138, MAE: 0.0913, RMSE: 0.1174, MAPE: 3.66%, R2: 0.6614, Time: 60.64s

XGBoost – MSE: 0.0043, MAE: 0.0495, RMSE: 0.0659, MAPE: 1.96%, R2: 0.8933, Time: 3.16s

CatBoost – MSE: 0.0043, MAE: 0.0494, RMSE: 0.0657, MAPE: 1.96%, R2: 0.8939, Time: 63.32s

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.560479

LightGBM – MSE: 0.0061, MAE: 0.0588, RMSE: 0.0781, MAPE: 2.33%, R2: 0.8501, Time: 5.47s

Results for target variable: ST117TA

Linear Regression – MSE: 0.0287, MAE: 0.1303, RMSE: 0.1695, MAPE: 5.15%, R2: 0.3011, Time: 0.36s

Decision Tree – MSE: 0.0037, MAE: 0.0444, RMSE: 0.0612, MAPE: 1.74%, R2: 0.9090, Time: 1.79s

Random Forest – MSE: 0.0037, MAE: 0.0443, RMSE: 0.0611, MAPE: 1.74%, R2: 0.9091, Time: 131.52s

k-NN – MSE: 0.0040, MAE: 0.0455, RMSE: 0.0631, MAPE: 1.78%, R2: 0.9031, Time: 7.94s

Gradient Boosting – MSE: 0.0129, MAE: 0.0872, RMSE: 0.1137, MAPE: 3.47%, R2: 0.6856, Time: 57.02s

XGBoost – MSE: 0.0039, MAE: 0.0457, RMSE: 0.0622, MAPE: 1.80%, R2: 0.9057, Time: 4.91s

CatBoost – MSE: 0.0038, MAE: 0.0457, RMSE: 0.0619, MAPE: 1.79%, R2: 0.9067, Time: 65.06s

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.582449

LightGBM – MSE: 0.0054, MAE: 0.0547, RMSE: 0.0738, MAPE: 2.16%, R2: 0.8674, Time: 5.08s

Results for target variable: ST40TA

Linear Regression – MSE: 0.0287, MAE: 0.1273, RMSE: 0.1693, MAPE: 4.95%, R2: 0.1904, Time: 0.22s

Decision Tree – MSE: 0.0043, MAE: 0.0468, RMSE: 0.0654, MAPE: 1.78%, R2: 0.8793, Time: 1.71s

Random Forest – MSE: 0.0043, MAE: 0.0468, RMSE: 0.0654, MAPE: 1.78%, R2: 0.8794, Time: 128.95s

k-NN – MSE: 0.0045, MAE: 0.0479, RMSE: 0.0674, MAPE: 1.83%, R2: 0.8716, Time: 6.84s

Gradient Boosting – MSE: 0.0145, MAE: 0.0901, RMSE: 0.1203, MAPE: 3.

49%, R2: 0.5910, Time: 59.60s
XGBoost – MSE: 0.0044, MAE: 0.0480, RMSE: 0.0662, MAPE: 1.83%, R2: 0.8761, Time: 3.28s
CatBoost – MSE: 0.0043, MAE: 0.0476, RMSE: 0.0655, MAPE: 1.82%, R2: 0.8788, Time: 63.93s
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.652801
LightGBM – MSE: 0.0060, MAE: 0.0567, RMSE: 0.0773, MAPE: 2.17%, R2: 0.8313, Time: 5.78s
Results for target variable: ST99TA
Linear Regression – MSE: 0.0277, MAE: 0.1256, RMSE: 0.1664, MAPE: 4.81%, R2: 0.1649, Time: 0.24s
Decision Tree – MSE: 0.0039, MAE: 0.0447, RMSE: 0.0625, MAPE: 1.70%, R2: 0.8823, Time: 1.77s
Random Forest – MSE: 0.0039, MAE: 0.0447, RMSE: 0.0624, MAPE: 1.70%, R2: 0.8824, Time: 130.89s
k-NN – MSE: 0.0042, MAE: 0.0457, RMSE: 0.0645, MAPE: 1.74%, R2: 0.8744, Time: 7.89s
Gradient Boosting – MSE: 0.0131, MAE: 0.0857, RMSE: 0.1147, MAPE: 3.29%, R2: 0.6033, Time: 60.36s
XGBoost – MSE: 0.0041, MAE: 0.0465, RMSE: 0.0637, MAPE: 1.77%, R2: 0.8776, Time: 3.20s
CatBoost – MSE: 0.0040, MAE: 0.0458, RMSE: 0.0629, MAPE: 1.74%, R2: 0.8808, Time: 64.09s
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.680324
LightGBM – MSE: 0.0056, MAE: 0.0552, RMSE: 0.0751, MAPE: 2.10%, R2: 0.8296, Time: 5.46s
Results for target variable: ST32TA
Linear Regression – MSE: 0.0302, MAE: 0.1328, RMSE: 0.1737, MAPE: 5.16%, R2: 0.1293, Time: 0.33s
Decision Tree – MSE: 0.0053, MAE: 0.0526, RMSE: 0.0727, MAPE: 2.01%, R2: 0.8476, Time: 2.02s
Random Forest – MSE: 0.0053, MAE: 0.0526, RMSE: 0.0726, MAPE: 2.01%, R2: 0.8478, Time: 131.13s
k-NN – MSE: 0.0056, MAE: 0.0536, RMSE: 0.0748, MAPE: 2.04%, R2: 0.8386, Time: 7.83s
Gradient Boosting – MSE: 0.0159, MAE: 0.0950, RMSE: 0.1261, MAPE: 3.67%, R2: 0.5414, Time: 59.00s
XGBoost – MSE: 0.0054, MAE: 0.0539, RMSE: 0.0733, MAPE: 2.06%, R2: 0.8450, Time: 3.29s
CatBoost – MSE: 0.0053, MAE: 0.0535, RMSE: 0.0728, MAPE: 2.04%, R2: 0.8473, Time: 63.67s
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.637400
LightGBM – MSE: 0.0071, MAE: 0.0630, RMSE: 0.0845, MAPE: 2.41%, R2: 0.7942, Time: 5.07s
Results for target variable: ST91TA
Linear Regression – MSE: 0.0282, MAE: 0.1288, RMSE: 0.1680, MAPE: 4.98%, R2: 0.1500, Time: 0.23s

Decision Tree – MSE: 0.0059, MAE: 0.0549, RMSE: 0.0768, MAPE: 2.11%, R2: 0.8221, Time: 2.12s
Random Forest – MSE: 0.0059, MAE: 0.0548, RMSE: 0.0768, MAPE: 2.10%, R2: 0.8223, Time: 129.31s
k-NN – MSE: 0.0063, MAE: 0.0562, RMSE: 0.0793, MAPE: 2.15%, R2: 0.8104, Time: 7.00s
Gradient Boosting – MSE: 0.0153, MAE: 0.0916, RMSE: 0.1239, MAPE: 3.53%, R2: 0.5377, Time: 60.16s
XGBoost – MSE: 0.0060, MAE: 0.0561, RMSE: 0.0772, MAPE: 2.15%, R2: 0.8204, Time: 3.11s
CatBoost – MSE: 0.0058, MAE: 0.0556, RMSE: 0.0765, MAPE: 2.13%, R2: 0.8238, Time: 65.51s
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.663719
LightGBM – MSE: 0.0078, MAE: 0.0651, RMSE: 0.0884, MAPE: 2.50%, R2: 0.7648, Time: 5.55s
Results for target variable: ST60TA
Linear Regression – MSE: 0.1342, MAE: 0.2705, RMSE: 0.3664, MAPE: 8.82%, R2: 0.0553, Time: 0.22s
Decision Tree – MSE: 0.0277, MAE: 0.1221, RMSE: 0.1663, MAPE: 3.83%, R2: 0.8054, Time: 1.94s
Random Forest – MSE: 0.0276, MAE: 0.1220, RMSE: 0.1662, MAPE: 3.82%, R2: 0.8056, Time: 148.71s
k-NN – MSE: 0.0293, MAE: 0.1250, RMSE: 0.1710, MAPE: 3.91%, R2: 0.7941, Time: 7.87s
Gradient Boosting – MSE: 0.0816, MAE: 0.2111, RMSE: 0.2856, MAPE: 6.82%, R2: 0.4259, Time: 62.15s
XGBoost – MSE: 0.0282, MAE: 0.1250, RMSE: 0.1678, MAPE: 3.92%, R2: 0.8018, Time: 3.29s
CatBoost – MSE: 0.0274, MAE: 0.1233, RMSE: 0.1656, MAPE: 3.87%, R2: 0.8069, Time: 63.37s
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 3.347149
LightGBM – MSE: 0.0358, MAE: 0.1413, RMSE: 0.1892, MAPE: 4.45%, R2: 0.7481, Time: 5.01s
Results for target variable: ST119TA
Linear Regression – MSE: 0.1285, MAE: 0.2619, RMSE: 0.3585, MAPE: 8.41%, R2: 0.1191, Time: 0.22s
Decision Tree – MSE: 0.0325, MAE: 0.1315, RMSE: 0.1803, MAPE: 3.97%, R2: 0.7773, Time: 1.91s
Random Forest – MSE: 0.0324, MAE: 0.1314, RMSE: 0.1801, MAPE: 3.96%, R2: 0.7777, Time: 146.80s
k-NN – MSE: 0.0343, MAE: 0.1341, RMSE: 0.1853, MAPE: 4.04%, R2: 0.7648, Time: 8.05s
Gradient Boosting – MSE: 0.0846, MAE: 0.2145, RMSE: 0.2909, MAPE: 6.80%, R2: 0.4198, Time: 60.47s
XGBoost – MSE: 0.0326, MAE: 0.1335, RMSE: 0.1804, MAPE: 4.04%, R2: 0.7769, Time: 3.34s
CatBoost – MSE: 0.0322, MAE: 0.1327, RMSE: 0.1793, MAPE: 4.02%, R2: 0.7795, Time: 63.83s
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, nu

mber of used features: 13

[LightGBM] [Info] Start training from score 3.431542

LightGBM – MSE: 0.0405, MAE: 0.1503, RMSE: 0.2013, MAPE: 4.60%, R2: 0.7222, Time: 4.92s

Results for target variable: ST62TA

Linear Regression – MSE: 0.0389, MAE: 0.1573, RMSE: 0.1974, MAPE: 7.10%, R2: 0.1289, Time: 0.22s

Decision Tree – MSE: 0.0067, MAE: 0.0615, RMSE: 0.0820, MAPE: 2.74%, R2: 0.8496, Time: 2.86s

Random Forest – MSE: 0.0067, MAE: 0.0615, RMSE: 0.0819, MAPE: 2.74%, R2: 0.8499, Time: 139.59s

k-NN – MSE: 0.0071, MAE: 0.0629, RMSE: 0.0844, MAPE: 2.81%, R2: 0.8406, Time: 6.96s

Gradient Boosting – MSE: 0.0218, MAE: 0.1147, RMSE: 0.1476, MAPE: 5.17%, R2: 0.5126, Time: 59.69s

XGBoost – MSE: 0.0069, MAE: 0.0630, RMSE: 0.0831, MAPE: 2.81%, R2: 0.8456, Time: 3.30s

CatBoost – MSE: 0.0068, MAE: 0.0624, RMSE: 0.0824, MAPE: 2.79%, R2: 0.8483, Time: 66.66s

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.270637

LightGBM – MSE: 0.0099, MAE: 0.0762, RMSE: 0.0995, MAPE: 3.41%, R2: 0.7786, Time: 4.98s

Results for target variable: ST121TA

Linear Regression – MSE: 0.0398, MAE: 0.1570, RMSE: 0.1995, MAPE: 6.94%, R2: 0.1548, Time: 0.22s

Decision Tree – MSE: 0.0081, MAE: 0.0665, RMSE: 0.0899, MAPE: 2.93%, R2: 0.8283, Time: 1.77s

Random Forest – MSE: 0.0081, MAE: 0.0664, RMSE: 0.0899, MAPE: 2.92%, R2: 0.8286, Time: 137.63s

k-NN – MSE: 0.0086, MAE: 0.0681, RMSE: 0.0926, MAPE: 3.00%, R2: 0.8179, Time: 7.79s

Gradient Boosting – MSE: 0.0238, MAE: 0.1168, RMSE: 0.1543, MAPE: 5.17%, R2: 0.4941, Time: 62.05s

XGBoost – MSE: 0.0082, MAE: 0.0677, RMSE: 0.0906, MAPE: 2.98%, R2: 0.8256, Time: 3.29s

CatBoost – MSE: 0.0082, MAE: 0.0677, RMSE: 0.0906, MAPE: 2.98%, R2: 0.8259, Time: 64.06s

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.302973

LightGBM – MSE: 0.0121, MAE: 0.0824, RMSE: 0.1099, MAPE: 3.64%, R2: 0.7435, Time: 5.01s

Voting Regressor Results for target variable: ST58TA

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.560479

Voting Regressor – MSE: 0.0055, MAE: 0.0571, RMSE: 0.0739, MAPE: 2.28%, R2: 0.8660, Time: 274.19s

Voting Regressor Results for target variable: ST117TA

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.582449
Voting Regressor – MSE: 0.0049, MAE: 0.0535, RMSE: 0.0701, MAPE: 2.11%, R2: 0.8803, Time: 263.63s

Voting Regressor Results for target variable: ST40TA

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.652801
Voting Regressor – MSE: 0.0055, MAE: 0.0551, RMSE: 0.0742, MAPE: 2.12%, R2: 0.8447, Time: 273.04s

Voting Regressor Results for target variable: ST99TA

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.680324
Voting Regressor – MSE: 0.0051, MAE: 0.0536, RMSE: 0.0712, MAPE: 2.05%, R2: 0.8470, Time: 273.05s

Voting Regressor Results for target variable: ST32TA

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.637400
Voting Regressor – MSE: 0.0066, MAE: 0.0611, RMSE: 0.0810, MAPE: 2.34%, R2: 0.8107, Time: 275.06s

Voting Regressor Results for target variable: ST91TA

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.663719
Voting Regressor – MSE: 0.0070, MAE: 0.0619, RMSE: 0.0838, MAPE: 2.39%, R2: 0.7884, Time: 263.58s

Voting Regressor Results for target variable: ST60TA

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 3.347149
Voting Regressor – MSE: 0.0333, MAE: 0.1368, RMSE: 0.1824, MAPE: 4.35%, R2: 0.7658, Time: 279.18s

Voting Regressor Results for target variable: ST119TA

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 3.431542
Voting Regressor – MSE: 0.0378, MAE: 0.1472, RMSE: 0.1944, MAPE: 4.56%, R2: 0.7409, Time: 304.36s

Voting Regressor Results for target variable: ST62TA

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.270637

Voting Regressor – MSE: 0.0085, MAE: 0.0727, RMSE: 0.0925, MAPE: 3.26%, R2: 0.8088, Time: 295.15s



Voting Regressor Results for target variable: ST121TA

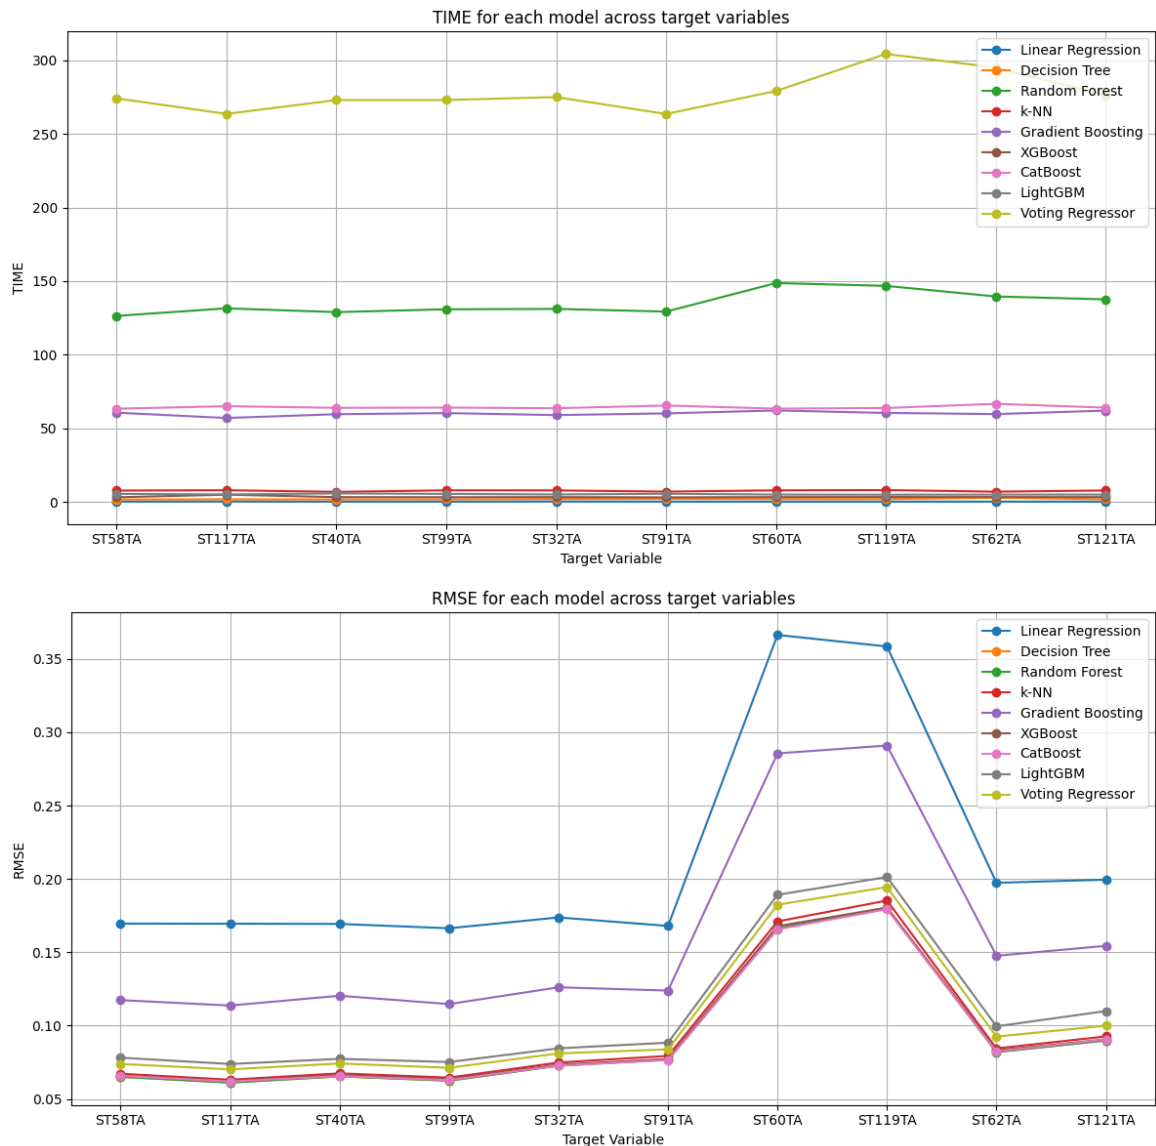
[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.302973

Voting Regressor – MSE: 0.0100, MAE: 0.0762, RMSE: 0.1000, MAPE: 3.37%, R2: 0.7878, Time: 276.31s





```
In [ ]: import pandas as pd
import numpy as np
import time
from itertools import combinations
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

import xgboost as xgb
from catboost import CatBoostRegressor
import lightgbm as lgb

# Assuming X and y are already defined DataFrames
# Example placeholder (remove if X and y are already defined)
# X = pd.read_csv("your_features.csv")
# y = pd.read_csv("your_targets.csv")

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize specified models
```

```

models = {
    'Random Forest': RandomForestRegressor(n_estimators=100),
    'Gradient Boosting': GradientBoostingRegressor(),
    'XGBoost': xgb.XGBRegressor(objective='reg:squarederror', verbose=0),
    'CatBoost': CatBoostRegressor(verbose=0),
    'LightGBM': lgb.LGBMRegressor(force_col_wise=True)
}

# Prepare to collect results for performance metrics
results = []

# Create combinations of specified models for Voting Regressor
model_names = list(models.keys())
for r in range(2, len(model_names) + 1):
    for combo in combinations(model_names, r):
        combo_name = ' + '.join(combo)
        combo_estimators = [(name, models[name]) for name in combo]
        voting_regressor = VotingRegressor(estimators=combo_estimators)

        for target_col in y_train.columns:
            print(f"Results for target variable: {target_col} with {combo_name}")
            start_time = time.time()
            voting_regressor.fit(X_train, y_train[target_col])
            y_pred = voting_regressor.predict(X_test)
            end_time = time.time()

            elapsed_time = end_time - start_time
            mse = mean_squared_error(y_test[target_col], y_pred)
            mae = mean_absolute_error(y_test[target_col], y_pred)
            rmse = np.sqrt(mse)
            mape = np.mean(np.abs((y_test[target_col] - y_pred) / y_test[target_col]))
            r2 = r2_score(y_test[target_col], y_pred)
            rse = np.sqrt(np.sum((y_test[target_col] - y_pred) ** 2) / len(y_test))

            print(f"Voting Regressor ({combo_name}) - MSE: {mse:.4f}")

            results.append({
                'model': f'Voting Regressor ({combo_name})',
                'target_col': target_col,
                'mse': mse,
                'mae': mae,
                'rmse': rmse,
                'mape': mape,
                'r2': r2,
                'rse': rse,
                'time': elapsed_time
            })

# Convert results to DataFrame for easy plotting
results_df = pd.DataFrame(results)

# Plotting results separately
for metric in ['mse', 'r2', 'time', 'rmse']:
    plt.figure(figsize=(12, 6))
    for name in results_df['model'].unique():
        model_results = results_df[results_df['model'] == name]

```



```
plt.plot(model_results['target_col'], model_results[metric])
plt.title(f'{metric.upper()} for each model combination')
plt.xlabel('Target Variable')
plt.ylabel(metric.upper())
plt.legend()
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```

Results for target variable: ST58TA with model combination: Random Forest + Gradient Boosting

Voting Regressor (Random Forest + Gradient Boosting) – MSE: 0.0065, MAE: 0.0627, RMSE: 0.0807, MAPE: 2.50%, R2: 0.8399, RSE: 0.0807, Time: 190.66s

Results for target variable: ST117TA with model combination: Random Forest + Gradient Boosting

Voting Regressor (Random Forest + Gradient Boosting) – MSE: 0.0059, MAE: 0.0589, RMSE: 0.0771, MAPE: 2.33%, R2: 0.8554, RSE: 0.0771, Time: 183.45s

Results for target variable: ST40TA with model combination: Random Forest + Gradient Boosting

Voting Regressor (Random Forest + Gradient Boosting) – MSE: 0.0067, MAE: 0.0612, RMSE: 0.0821, MAPE: 2.36%, R2: 0.8098, RSE: 0.0821, Time: 178.92s

Results for target variable: ST99TA with model combination: Random Forest + Gradient Boosting

Voting Regressor (Random Forest + Gradient Boosting) – MSE: 0.0061, MAE: 0.0589, RMSE: 0.0784, MAPE: 2.25%, R2: 0.8147, RSE: 0.0784, Time: 183.06s

Results for target variable: ST32TA with model combination: Random Forest + Gradient Boosting

Voting Regressor (Random Forest + Gradient Boosting) – MSE: 0.0079, MAE: 0.0668, RMSE: 0.0886, MAPE: 2.57%, R2: 0.7736, RSE: 0.0886, Time: 189.33s

Results for target variable: ST91TA with model combination: Random Forest + Gradient Boosting

Voting Regressor (Random Forest + Gradient Boosting) – MSE: 0.0081, MAE: 0.0670, RMSE: 0.0902, MAPE: 2.58%, R2: 0.7548, RSE: 0.0902, Time: 176.95s

Results for target variable: ST60TA with model combination: Random Forest + Gradient Boosting

Voting Regressor (Random Forest + Gradient Boosting) – MSE: 0.0404, MAE: 0.1506, RMSE: 0.2010, MAPE: 4.81%, R2: 0.7156, RSE: 0.2010, Time: 199.63s

Results for target variable: ST119TA with model combination: Random Forest + Gradient Boosting

Voting Regressor (Random Forest + Gradient Boosting) – MSE: 0.0449, MAE: 0.1600, RMSE: 0.2118, MAPE: 4.99%, R2: 0.6926, RSE: 0.2118, Time: 195.57s

Results for target variable: ST62TA with model combination: Random Forest + Gradient Boosting

Voting Regressor (Random Forest + Gradient Boosting) – MSE: 0.0103, MAE: 0.0799, RMSE: 0.1017, MAPE: 3.59%, R2: 0.7689, RSE: 0.1017, Time: 188.50s

Results for target variable: ST121TA with model combination: Random

Forest + Gradient Boosting

Voting Regressor (Random Forest + Gradient Boosting) – MSE: 0.0118, MAE: 0.0830, RMSE: 0.1088, MAPE: 3.67%, R2: 0.7488, RSE: 0.1088, Time: 201.56s

Results for target variable: ST58TA with model combination: Random Forest + XGBoost

Voting Regressor (Random Forest + XGBoost) – MSE: 0.0042, MAE: 0.0484, RMSE: 0.0647, MAPE: 1.91%, R2: 0.8972, RSE: 0.0647, Time: 142.71s

Results for target variable: ST117TA with model combination: Random Forest + XGBoost

Voting Regressor (Random Forest + XGBoost) – MSE: 0.0037, MAE: 0.0445, RMSE: 0.0609, MAPE: 1.75%, R2: 0.9098, RSE: 0.0609, Time: 134.70s

Results for target variable: ST40TA with model combination: Random Forest + XGBoost

Voting Regressor (Random Forest + XGBoost) – MSE: 0.0042, MAE: 0.0468, RMSE: 0.0650, MAPE: 1.79%, R2: 0.8808, RSE: 0.0650, Time: 129.19s

Results for target variable: ST99TA with model combination: Random Forest + XGBoost

Voting Regressor (Random Forest + XGBoost) – MSE: 0.0039, MAE: 0.0450, RMSE: 0.0622, MAPE: 1.71%, R2: 0.8831, RSE: 0.0622, Time: 132.74s

Results for target variable: ST32TA with model combination: Random Forest + XGBoost

Voting Regressor (Random Forest + XGBoost) – MSE: 0.0052, MAE: 0.0527, RMSE: 0.0722, MAPE: 2.01%, R2: 0.8497, RSE: 0.0722, Time: 137.44s

Results for target variable: ST91TA with model combination: Random Forest + XGBoost

Voting Regressor (Random Forest + XGBoost) – MSE: 0.0058, MAE: 0.0549, RMSE: 0.0762, MAPE: 2.11%, R2: 0.8248, RSE: 0.0762, Time: 125.44s

Results for target variable: ST60TA with model combination: Random Forest + XGBoost

Voting Regressor (Random Forest + XGBoost) – MSE: 0.0273, MAE: 0.1222, RMSE: 0.1651, MAPE: 3.83%, R2: 0.8081, RSE: 0.1651, Time: 149.96s

Results for target variable: ST119TA with model combination: Random Forest + XGBoost

Voting Regressor (Random Forest + XGBoost) – MSE: 0.0319, MAE: 0.1313, RMSE: 0.1786, MAPE: 3.97%, R2: 0.7813, RSE: 0.1786, Time: 144.41s

Results for target variable: ST62TA with model combination: Random Forest + XGBoost

Voting Regressor (Random Forest + XGBoost) – MSE: 0.0066, MAE: 0.0615, RMSE: 0.0815, MAPE: 2.75%, R2: 0.8515, RSE: 0.0815, Time: 141.59s

Results for target variable: ST121TA with model combination: Random Forest + XGBoost

Voting Regressor (Random Forest + XGBoost) – MSE: 0.0080, MAE: 0.0664, RMSE: 0.0892, MAPE: 2.92%, R2: 0.8309, RSE: 0.0892, Time: 143.42s

Results for target variable: ST58TA with model combination: Random Forest + CatBoost

Voting Regressor (Random Forest + CatBoost) – MSE: 0.0042, MAE: 0.0483, RMSE: 0.0646, MAPE: 1.91%, R2: 0.8974, RSE: 0.0646, Time: 198.08s

Results for target variable: ST117TA with model combination: Random Forest + CatBoost

Voting Regressor (Random Forest + CatBoost) – MSE: 0.0037, MAE: 0.0445, RMSE: 0.0608, MAPE: 1.74%, R2: 0.9101, RSE: 0.0608, Time: 193.05s

Results for target variable: ST40TA with model combination: Random Forest + CatBoost

Voting Regressor (Random Forest + CatBoost) – MSE: 0.0042, MAE: 0.0467, RMSE: 0.0648, MAPE: 1.78%, R2: 0.8815, RSE: 0.0648, Time: 189.12 s

Results for target variable: ST99TA with model combination: Random Forest + CatBoost

Voting Regressor (Random Forest + CatBoost) – MSE: 0.0038, MAE: 0.0448, RMSE: 0.0620, MAPE: 1.70%, R2: 0.8840, RSE: 0.0620, Time: 194.50 s

Results for target variable: ST32TA with model combination: Random Forest + CatBoost

Voting Regressor (Random Forest + CatBoost) – MSE: 0.0052, MAE: 0.0526, RMSE: 0.0720, MAPE: 2.01%, R2: 0.8503, RSE: 0.0720, Time: 200.22 s

Results for target variable: ST91TA with model combination: Random Forest + CatBoost

Voting Regressor (Random Forest + CatBoost) – MSE: 0.0058, MAE: 0.0547, RMSE: 0.0760, MAPE: 2.10%, R2: 0.8258, RSE: 0.0760, Time: 184.25 s

Results for target variable: ST60TA with model combination: Random Forest + CatBoost

Voting Regressor (Random Forest + CatBoost) – MSE: 0.0270, MAE: 0.1216, RMSE: 0.1645, MAPE: 3.81%, R2: 0.8096, RSE: 0.1645, Time: 210.80 s

Results for target variable: ST119TA with model combination: Random Forest + CatBoost

Voting Regressor (Random Forest + CatBoost) – MSE: 0.0318, MAE: 0.1310, RMSE: 0.1783, MAPE: 3.96%, R2: 0.7821, RSE: 0.1783, Time: 204.76 s

Results for target variable: ST62TA with model combination: Random Forest + CatBoost

Voting Regressor (Random Forest + CatBoost) – MSE: 0.0066, MAE: 0.0614, RMSE: 0.0813, MAPE: 2.74%, R2: 0.8523, RSE: 0.0813, Time: 203.02 s

Results for target variable: ST121TA with model combination: Random Forest + CatBoost

Voting Regressor (Random Forest + CatBoost) – MSE: 0.0080, MAE: 0.0664, RMSE: 0.0892, MAPE: 2.92%, R2: 0.8310, RSE: 0.0892, Time: 203.56 s

Results for target variable: ST58TA with model combination: Random Forest + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.560479

Voting Regressor (Random Forest + LightGBM) – MSE: 0.0046, MAE: 0.0512, RMSE: 0.0679, MAPE: 2.03%, R2: 0.8867, RSE: 0.0679, Time: 139.21 s

Results for target variable: ST117TA with model combination: Random Forest + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.582449

Voting Regressor (Random Forest + LightGBM) – MSE: 0.0041, MAE: 0.0473, RMSE: 0.0639, MAPE: 1.86%, R2: 0.9005, RSE: 0.0639, Time: 132.87 s

Results for target variable: ST40TA with model combination: Random Forest + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.652801

Voting Regressor (Random Forest + LightGBM) – MSE: 0.0046, MAE: 0.0495, RMSE: 0.0679, MAPE: 1.89%, R2: 0.8698, RSE: 0.0679, Time: 133.37 s

Results for target variable: ST99TA with model combination: Random Forest + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.680324

Voting Regressor (Random Forest + LightGBM) – MSE: 0.0043, MAE: 0.0478, RMSE: 0.0653, MAPE: 1.82%, R2: 0.8713, RSE: 0.0653, Time: 137.09 s

Results for target variable: ST32TA with model combination: Random Forest + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.637400

Voting Regressor (Random Forest + LightGBM) – MSE: 0.0057, MAE: 0.0556, RMSE: 0.0752, MAPE: 2.12%, R2: 0.8370, RSE: 0.0752, Time: 141.86 s

Results for target variable: ST91TA with model combination: Random Forest + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.663719

Voting Regressor (Random Forest + LightGBM) – MSE: 0.0063, MAE: 0.0577, RMSE: 0.0791, MAPE: 2.22%, R2: 0.8114, RSE: 0.0791, Time: 128.74 s

Results for target variable: ST60TA with model combination: Random Forest + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 3.347149

Voting Regressor (Random Forest + LightGBM) – MSE: 0.0291, MAE: 0.1271, RMSE: 0.1706, MAPE: 4.00%, R2: 0.7952, RSE: 0.1706, Time: 155.02 s

Results for target variable: ST119TA with model combination: Random Forest + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 3.431542

Voting Regressor (Random Forest + LightGBM) – MSE: 0.0338, MAE: 0.1365, RMSE: 0.1840, MAPE: 4.15%, R2: 0.7680, RSE: 0.1840, Time: 144.35 s

Results for target variable: ST62TA with model combination: Random Forest + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.270637
Voting Regressor (Random Forest + LightGBM) – MSE: 0.0074, MAE: 0.0658, RMSE: 0.0859, MAPE: 2.94%, R2: 0.8351, RSE: 0.0859, Time: 144.53s
Results for target variable: ST121TA with model combination: Random Forest + LightGBM
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.302973
Voting Regressor (Random Forest + LightGBM) – MSE: 0.0089, MAE: 0.0711, RMSE: 0.0944, MAPE: 3.14%, R2: 0.8109, RSE: 0.0944, Time: 139.32s
Results for target variable: ST58TA with model combination: Gradient Boosting + XGBoost
Voting Regressor (Gradient Boosting + XGBoost) – MSE: 0.0070, MAE: 0.0648, RMSE: 0.0834, MAPE: 2.59%, R2: 0.8290, RSE: 0.0834, Time: 63.17s
Results for target variable: ST117TA with model combination: Gradient Boosting + XGBoost
Voting Regressor (Gradient Boosting + XGBoost) – MSE: 0.0064, MAE: 0.0609, RMSE: 0.0798, MAPE: 2.41%, R2: 0.8452, RSE: 0.0798, Time: 64.83s
Results for target variable: ST40TA with model combination: Gradient Boosting + XGBoost
Voting Regressor (Gradient Boosting + XGBoost) – MSE: 0.0072, MAE: 0.0632, RMSE: 0.0847, MAPE: 2.44%, R2: 0.7976, RSE: 0.0847, Time: 65.78s
Results for target variable: ST99TA with model combination: Gradient Boosting + XGBoost
Voting Regressor (Gradient Boosting + XGBoost) – MSE: 0.0066, MAE: 0.0611, RMSE: 0.0812, MAPE: 2.34%, R2: 0.8013, RSE: 0.0812, Time: 63.42s
Results for target variable: ST32TA with model combination: Gradient Boosting + XGBoost
Voting Regressor (Gradient Boosting + XGBoost) – MSE: 0.0083, MAE: 0.0688, RMSE: 0.0911, MAPE: 2.64%, R2: 0.7605, RSE: 0.0911, Time: 64.33s
Results for target variable: ST91TA with model combination: Gradient Boosting + XGBoost
Voting Regressor (Gradient Boosting + XGBoost) – MSE: 0.0086, MAE: 0.0690, RMSE: 0.0926, MAPE: 2.66%, R2: 0.7414, RSE: 0.0926, Time: 63.02s
Results for target variable: ST60TA with model combination: Gradient Boosting + XGBoost
Voting Regressor (Gradient Boosting + XGBoost) – MSE: 0.0428, MAE: 0.1547, RMSE: 0.2068, MAPE: 4.94%, R2: 0.6990, RSE: 0.2068, Time: 62.94s
Results for target variable: ST119TA with model combination: Gradient Boosting + XGBoost
Voting Regressor (Gradient Boosting + XGBoost) – MSE: 0.0470, MAE: 0.1637, RMSE: 0.2168, MAPE: 5.11%, R2: 0.6780, RSE: 0.2168, Time: 63.02s

Results for target variable: ST62TA with model combination: Gradient Boosting + XGBoost

Voting Regressor (Gradient Boosting + XGBoost) – MSE: 0.0111, MAE: 0.0824, RMSE: 0.1053, MAPE: 3.70%, R2: 0.7522, RSE: 0.1053, Time: 6 3.30s

Results for target variable: ST121TA with model combination: Gradient Boosting + XGBoost

Voting Regressor (Gradient Boosting + XGBoost) – MSE: 0.0126, MAE: 0.0855, RMSE: 0.1124, MAPE: 3.78%, R2: 0.7316, RSE: 0.1124, Time: 6 2.47s

Results for target variable: ST58TA with model combination: Gradient Boosting + CatBoost

Voting Regressor (Gradient Boosting + CatBoost) – MSE: 0.0070, MAE: 0.0648, RMSE: 0.0834, MAPE: 2.59%, R2: 0.8289, RSE: 0.0834, Time: 12 0.27s

Results for target variable: ST117TA with model combination: Gradient Boosting + CatBoost

Voting Regressor (Gradient Boosting + CatBoost) – MSE: 0.0063, MAE: 0.0609, RMSE: 0.0796, MAPE: 2.41%, R2: 0.8457, RSE: 0.0796, Time: 12 1.30s

Results for target variable: ST40TA with model combination: Gradient Boosting + CatBoost

Voting Regressor (Gradient Boosting + CatBoost) – MSE: 0.0071, MAE: 0.0629, RMSE: 0.0842, MAPE: 2.42%, R2: 0.7998, RSE: 0.0842, Time: 12 1.10s

Results for target variable: ST99TA with model combination: Gradient Boosting + CatBoost

Voting Regressor (Gradient Boosting + CatBoost) – MSE: 0.0065, MAE: 0.0606, RMSE: 0.0806, MAPE: 2.32%, R2: 0.8042, RSE: 0.0806, Time: 12 2.42s

Results for target variable: ST32TA with model combination: Gradient Boosting + CatBoost

Voting Regressor (Gradient Boosting + CatBoost) – MSE: 0.0082, MAE: 0.0685, RMSE: 0.0908, MAPE: 2.64%, R2: 0.7624, RSE: 0.0908, Time: 12 3.30s

Results for target variable: ST91TA with model combination: Gradient Boosting + CatBoost

Voting Regressor (Gradient Boosting + CatBoost) – MSE: 0.0085, MAE: 0.0685, RMSE: 0.0921, MAPE: 2.64%, R2: 0.7445, RSE: 0.0921, Time: 12 2.13s

Results for target variable: ST60TA with model combination: Gradient Boosting + CatBoost

Voting Regressor (Gradient Boosting + CatBoost) – MSE: 0.0422, MAE: 0.1539, RMSE: 0.2054, MAPE: 4.92%, R2: 0.7032, RSE: 0.2054, Time: 11 9.01s

Results for target variable: ST119TA with model combination: Gradient Boosting + CatBoost

Voting Regressor (Gradient Boosting + CatBoost) – MSE: 0.0467, MAE: 0.1632, RMSE: 0.2160, MAPE: 5.09%, R2: 0.6802, RSE: 0.2160, Time: 12 1.34s

Results for target variable: ST62TA with model combination: Gradient Boosting + CatBoost

Voting Regressor (Gradient Boosting + CatBoost) – MSE: 0.0110, MAE: 0.0821, RMSE: 0.1048, MAPE: 3.69%, R2: 0.7545, RSE: 0.1048, Time: 12 0.89s

Results for target variable: ST121TA with model combination: Gradient

t Boosting + CatBoost

Voting Regressor (Gradient Boosting + CatBoost) – MSE: 0.0126, MAE: 0.0855, RMSE: 0.1125, MAPE: 3.78%, R2: 0.7314, RSE: 0.1125, Time: 12 2.43s

Results for target variable: ST58TA with model combination: Gradient Boosting + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.560479

Voting Regressor (Gradient Boosting + LightGBM) – MSE: 0.0088, MAE: 0.0726, RMSE: 0.0940, MAPE: 2.90%, R2: 0.7831, RSE: 0.0940, Time: 6 3.30s

Results for target variable: ST117TA with model combination: Gradient Boosting + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.582449

Voting Regressor (Gradient Boosting + LightGBM) – MSE: 0.0080, MAE: 0.0685, RMSE: 0.0897, MAPE: 2.71%, R2: 0.8042, RSE: 0.0897, Time: 6 3.96s

Results for target variable: ST40TA with model combination: Gradient Boosting + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.652801

Voting Regressor (Gradient Boosting + LightGBM) – MSE: 0.0089, MAE: 0.0707, RMSE: 0.0945, MAPE: 2.73%, R2: 0.7477, RSE: 0.0945, Time: 6 4.78s

Results for target variable: ST99TA with model combination: Gradient Boosting + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.680324

Voting Regressor (Gradient Boosting + LightGBM) – MSE: 0.0083, MAE: 0.0681, RMSE: 0.0909, MAPE: 2.60%, R2: 0.7507, RSE: 0.0909, Time: 6 3.10s

Results for target variable: ST32TA with model combination: Gradient Boosting + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.637400

Voting Regressor (Gradient Boosting + LightGBM) – MSE: 0.0102, MAE: 0.0763, RMSE: 0.1011, MAPE: 2.94%, R2: 0.7050, RSE: 0.1011, Time: 6 4.78s

Results for target variable: ST91TA with model combination: Gradient Boosting + LightGBM

[LightGBM] [Info] Total Bins 1102

[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13

[LightGBM] [Info] Start training from score 2.663719

Voting Regressor (Gradient Boosting + LightGBM) – MSE: 0.0105, MAE:

0.0762, RMSE: 0.1027, MAPE: 2.94%, R2: 0.6821, RSE: 0.1027, Time: 6 3.67s
Results for target variable: ST60TA with model combination: Gradient Boosting + LightGBM
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 3.347149
Voting Regressor (Gradient Boosting + LightGBM) – MSE: 0.0515, MAE: 0.1697, RMSE: 0.2269, MAPE: 5.42%, R2: 0.6377, RSE: 0.2269, Time: 6 3.97s
Results for target variable: ST119TA with model combination: Gradient Boosting + LightGBM
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 3.431542
Voting Regressor (Gradient Boosting + LightGBM) – MSE: 0.0560, MAE: 0.1780, RMSE: 0.2366, MAPE: 5.57%, R2: 0.6163, RSE: 0.2366, Time: 6 3.60s
Results for target variable: ST62TA with model combination: Gradient Boosting + LightGBM
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.270637
Voting Regressor (Gradient Boosting + LightGBM) – MSE: 0.0143, MAE: 0.0929, RMSE: 0.1195, MAPE: 4.18%, R2: 0.6808, RSE: 0.1195, Time: 6 4.19s
Results for target variable: ST121TA with model combination: Gradient Boosting + LightGBM
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.302973
Voting Regressor (Gradient Boosting + LightGBM) – MSE: 0.0165, MAE: 0.0972, RMSE: 0.1285, MAPE: 4.30%, R2: 0.6492, RSE: 0.1285, Time: 6 4.29s
Results for target variable: ST58TA with model combination: XGBoost + CatBoost
Voting Regressor (XGBoost + CatBoost) – MSE: 0.0043, MAE: 0.0492, RMSE: 0.0654, MAPE: 1.95%, R2: 0.8948, RSE: 0.0654, Time: 65.99s
Results for target variable: ST117TA with model combination: XGBoost + CatBoost
Voting Regressor (XGBoost + CatBoost) – MSE: 0.0038, MAE: 0.0453, RMSE: 0.0617, MAPE: 1.78%, R2: 0.9075, RSE: 0.0617, Time: 65.98s
Results for target variable: ST40TA with model combination: XGBoost + CatBoost
Voting Regressor (XGBoost + CatBoost) – MSE: 0.0043, MAE: 0.0475, RMSE: 0.0655, MAPE: 1.81%, R2: 0.8788, RSE: 0.0655, Time: 68.55s
Results for target variable: ST99TA with model combination: XGBoost + CatBoost
Voting Regressor (XGBoost + CatBoost) – MSE: 0.0040, MAE: 0.0459, RMSE: 0.0629, MAPE: 1.74%, R2: 0.8806, RSE: 0.0629, Time: 66.12s
Results for target variable: ST32TA with model combination: XGBoost + CatBoost

Voting Regressor (XGBoost + CatBoost) – MSE: 0.0053, MAE: 0.0534, RMSE: 0.0727, MAPE: 2.04%, R2: 0.8476, RSE: 0.0727, Time: 65.59s
Results for target variable: ST91TA with model combination: XGBoost + CatBoost

Voting Regressor (XGBoost + CatBoost) – MSE: 0.0059, MAE: 0.0555, RMSE: 0.0765, MAPE: 2.13%, R2: 0.8237, RSE: 0.0765, Time: 65.78s
Results for target variable: ST60TA with model combination: XGBoost + CatBoost

Voting Regressor (XGBoost + CatBoost) – MSE: 0.0275, MAE: 0.1234, RMSE: 0.1659, MAPE: 3.87%, R2: 0.8064, RSE: 0.1659, Time: 68.89s
Results for target variable: ST119TA with model combination: XGBoost + CatBoost

Voting Regressor (XGBoost + CatBoost) – MSE: 0.0321, MAE: 0.1325, RMSE: 0.1792, MAPE: 4.01%, R2: 0.7800, RSE: 0.1792, Time: 65.66s
Results for target variable: ST62TA with model combination: XGBoost + CatBoost

Voting Regressor (XGBoost + CatBoost) – MSE: 0.0068, MAE: 0.0624, RMSE: 0.0824, MAPE: 2.79%, R2: 0.8483, RSE: 0.0824, Time: 65.80s
Results for target variable: ST121TA with model combination: XGBoost + CatBoost

Voting Regressor (XGBoost + CatBoost) – MSE: 0.0081, MAE: 0.0673, RMSE: 0.0902, MAPE: 2.97%, R2: 0.8273, RSE: 0.0902, Time: 67.39s
Results for target variable: ST58TA with model combination: XGBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.560479

Voting Regressor (XGBoost + LightGBM) – MSE: 0.0049, MAE: 0.0528, RMSE: 0.0700, MAPE: 2.09%, R2: 0.8795, RSE: 0.0700, Time: 8.42s
Results for target variable: ST117TA with model combination: XGBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.582449

Voting Regressor (XGBoost + LightGBM) – MSE: 0.0044, MAE: 0.0489, RMSE: 0.0661, MAPE: 1.92%, R2: 0.8938, RSE: 0.0661, Time: 8.99s
Results for target variable: ST40TA with model combination: XGBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.652801

Voting Regressor (XGBoost + LightGBM) – MSE: 0.0049, MAE: 0.0510, RMSE: 0.0699, MAPE: 1.95%, R2: 0.8619, RSE: 0.0699, Time: 9.61s
Results for target variable: ST99TA with model combination: XGBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.680324

Voting Regressor (XGBoost + LightGBM) – MSE: 0.0046, MAE: 0.0496, RMSE: 0.0675, MAPE: 1.89%, R2: 0.8624, RSE: 0.0675, Time: 8.97s
Results for target variable: ST32TA with model combination: XGBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.637400
Voting Regressor (XGBoost + LightGBM) – MSE: 0.0059, MAE: 0.0572, RMSE: 0.0771, MAPE: 2.18%, R2: 0.8287, RSE: 0.0771, Time: 10.38s
Results for target variable: ST91TA with model combination: XGBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.663719
Voting Regressor (XGBoost + LightGBM) – MSE: 0.0065, MAE: 0.0593, RMSE: 0.0808, MAPE: 2.28%, R2: 0.8031, RSE: 0.0808, Time: 8.22s
Results for target variable: ST60TA with model combination: XGBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 3.347149
Voting Regressor (XGBoost + LightGBM) – MSE: 0.0304, MAE: 0.1302, RMSE: 0.1745, MAPE: 4.10%, R2: 0.7857, RSE: 0.1745, Time: 9.77s
Results for target variable: ST119TA with model combination: XGBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 3.431542
Voting Regressor (XGBoost + LightGBM) – MSE: 0.0350, MAE: 0.1392, RMSE: 0.1871, MAPE: 4.24%, R2: 0.7599, RSE: 0.1871, Time: 8.82s
Results for target variable: ST62TA with model combination: XGBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.270637
Voting Regressor (XGBoost + LightGBM) – MSE: 0.0079, MAE: 0.0678, RMSE: 0.0886, MAPE: 3.03%, R2: 0.8242, RSE: 0.0886, Time: 8.08s
Results for target variable: ST121TA with model combination: XGBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.302973
Voting Regressor (XGBoost + LightGBM) – MSE: 0.0094, MAE: 0.0730, RMSE: 0.0972, MAPE: 3.23%, R2: 0.7993, RSE: 0.0972, Time: 9.49s
Results for target variable: ST58TA with model combination: CatBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.560479
Voting Regressor (CatBoost + LightGBM) – MSE: 0.0049, MAE: 0.0528, RMSE: 0.0700, MAPE: 2.09%, R2: 0.8797, RSE: 0.0700, Time: 67.61s
Results for target variable: ST117TA with model combination: CatBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.582449
Voting Regressor (CatBoost + LightGBM) – MSE: 0.0043, MAE: 0.0488, R MSE: 0.0659, MAPE: 1.92%, R2: 0.8944, RSE: 0.0659, Time: 67.20s
Results for target variable: ST40TA with model combination: CatBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.652801
Voting Regressor (CatBoost + LightGBM) – MSE: 0.0048, MAE: 0.0507, R MSE: 0.0694, MAPE: 1.94%, R2: 0.8638, RSE: 0.0694, Time: 67.04s
Results for target variable: ST99TA with model combination: CatBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.680324
Voting Regressor (CatBoost + LightGBM) – MSE: 0.0045, MAE: 0.0492, R MSE: 0.0670, MAPE: 1.87%, R2: 0.8646, RSE: 0.0670, Time: 70.50s
Results for target variable: ST32TA with model combination: CatBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.637400
Voting Regressor (CatBoost + LightGBM) – MSE: 0.0059, MAE: 0.0570, R MSE: 0.0767, MAPE: 2.18%, R2: 0.8302, RSE: 0.0767, Time: 67.73s
Results for target variable: ST91TA with model combination: CatBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.663719
Voting Regressor (CatBoost + LightGBM) – MSE: 0.0065, MAE: 0.0590, R MSE: 0.0804, MAPE: 2.27%, R2: 0.8054, RSE: 0.0804, Time: 67.82s
Results for target variable: ST60TA with model combination: CatBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 3.347149
Voting Regressor (CatBoost + LightGBM) – MSE: 0.0301, MAE: 0.1296, R MSE: 0.1734, MAPE: 4.07%, R2: 0.7883, RSE: 0.1734, Time: 67.24s
Results for target variable: ST119TA with model combination: CatBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 3.431542
Voting Regressor (CatBoost + LightGBM) – MSE: 0.0348, MAE: 0.1388, R MSE: 0.1866, MAPE: 4.23%, R2: 0.7614, RSE: 0.1866, Time: 69.45s
Results for target variable: ST62TA with model combination: CatBoost + LightGBM

[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.270637
Voting Regressor (CatBoost + LightGBM) – MSE: 0.0078, MAE: 0.0675, RMSE: 0.0882, MAPE: 3.02%, R2: 0.8260, RSE: 0.0882, Time: 67.47s
Results for target variable: ST121TA with model combination: CatBoost + LightGBM
[LightGBM] [Info] Total Bins 1102
[LightGBM] [Info] Number of data points in the train set: 430292, number of used features: 13
[LightGBM] [Info] Start training from score 2.302973
Voting Regressor (CatBoost + LightGBM) – MSE: 0.0095, MAE: 0.0731, RMSE: 0.0973, MAPE: 3.23%, R2: 0.7991, RSE: 0.0973, Time: 68.01s
Results for target variable: ST58TA with model combination: Random Forest + Gradient Boosting + XGBoost
Voting Regressor (Random Forest + Gradient Boosting + XGBoost) – MSE: 0.0054, MAE: 0.0564, RMSE: 0.0733, MAPE: 2.25%, R2: 0.8681, RSE: 0.0733, Time: 193.23s
Results for target variable: ST117TA with model combination: Random Forest + Gradient Boosting + XGBoost
Voting Regressor (Random Forest + Gradient Boosting + XGBoost) – MSE: 0.0048, MAE: 0.0525, RMSE: 0.0695, MAPE: 2.07%, R2: 0.8824, RSE: 0.0695, Time: 181.66s
Results for target variable: ST40TA with model combination: Random Forest + Gradient Boosting + XGBoost
Voting Regressor (Random Forest + Gradient Boosting + XGBoost) – MSE: 0.0055, MAE: 0.0548, RMSE: 0.0741, MAPE: 2.11%, R2: 0.8451, RSE: 0.0741, Time: 180.76s
Results for target variable: ST99TA with model combination: Random Forest + Gradient Boosting + XGBoost
Voting Regressor (Random Forest + Gradient Boosting + XGBoost) – MSE: 0.0050, MAE: 0.0530, RMSE: 0.0709, MAPE: 2.02%, R2: 0.8481, RSE: 0.0709, Time: 182.56s
Results for target variable: ST32TA with model combination: Random Forest + Gradient Boosting + XGBoost
Voting Regressor (Random Forest + Gradient Boosting + XGBoost) – MSE: 0.0065, MAE: 0.0608, RMSE: 0.0809, MAPE: 2.33%, R2: 0.8112, RSE: 0.0809, Time: 184.23s
Results for target variable: ST91TA with model combination: Random Forest + Gradient Boosting + XGBoost
Voting Regressor (Random Forest + Gradient Boosting + XGBoost) – MSE: 0.0070, MAE: 0.0618, RMSE: 0.0836, MAPE: 2.38%, R2: 0.7894, RSE: 0.0836, Time: 181.68s
Results for target variable: ST60TA with model combination: Random Forest + Gradient Boosting + XGBoost
Voting Regressor (Random Forest + Gradient Boosting + XGBoost) – MSE: 0.0338, MAE: 0.1380, RMSE: 0.1839, MAPE: 4.39%, R2: 0.7619, RSE: 0.1839, Time: 207.83s
Results for target variable: ST119TA with model combination: Random Forest + Gradient Boosting + XGBoost
Voting Regressor (Random Forest + Gradient Boosting + XGBoost) – MSE: 0.0384, MAE: 0.1477, RMSE: 0.1959, MAPE: 4.57%, R2: 0.7371, RSE: 0.1959, Time: 198.86s
Results for target variable: ST62TA with model combination: Random Forest + Gradient Boosting + XGBoost

Voting Regressor (Random Forest + Gradient Boosting + XGBoost) – MSE: 0.0085, MAE: 0.0720, RMSE: 0.0924, MAPE: 3.23%, R2: 0.8091, RSE: 0.0924, Time: 193.96s
 Results for target variable: ST121TA with model combination: Random Forest + Gradient Boosting + XGBoost
 Voting Regressor (Random Forest + Gradient Boosting + XGBoost) – MSE: 0.0100, MAE: 0.0758, RMSE: 0.0998, MAPE: 3.35%, R2: 0.7886, RSE: 0.0998, Time: 197.75s
 Results for target variable: ST58TA with model combination: Random Forest + Gradient Boosting + CatBoost

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt

# Corrected data format (shortened for demonstration – extend as needed)
data = {
    'Output': ['ST58TA', 'ST117TA', 'ST40TA', 'ST99TA', 'ST32TA'] * 5,
    'Model': ['RF + LGB + XGB'] * 5 + ['GB + XGB + CB'] * 5 + ['RF + LGB + XGB + CB'] * 5,
    'MSE': [0.0076, 0.0065, 0.0092, 0.0081, 0.0095,
            0.0076, 0.0065, 0.0092, 0.0080, 0.0094,
            0.0047, 0.0042, 0.0066, 0.0052, 0.0064],
    'RMSE': [0.0875, 0.0808, 0.0960, 0.0901, 0.0976,
            0.0875, 0.0808, 0.0960, 0.0896, 0.0971,
            0.0683, 0.0644, 0.0815, 0.0725, 0.0802],
    'R2': [0.8334, 0.8517, 0.8084, 0.8233, 0.8085,
            0.8336, 0.8519, 0.8086, 0.8245, 0.8096,
            0.8970, 0.9031, 0.8563, 0.8924, 0.8711],
    'Time(s)': [61.21, 57.84, 67.43, 63.67, 61.98,
                94.78, 90.52, 98.95, 96.55, 92.98,
                86.84, 83.23, 92.78, 90.97, 85.53]
}

df = pd.DataFrame(data)

# Plotting helper function
def plot_metric(metric, ylabel):
    plt.figure(figsize=(10, 6))
    for model in df['Model'].unique():
        subset = df[df['Model'] == model]
        plt.plot(subset['Output'], subset[metric], marker='o', linecolor='red')
    plt.title(f'{ylabel} for each Model')
    plt.xlabel('Output')
    plt.ylabel(ylabel)
    plt.xticks(rotation=45)
    plt.legend()
    plt.tight_layout()
    plt.show()

# Plot all metrics
plot_metric('MSE', 'MSE')
plot_metric('RMSE', 'RMSE')
plot_metric('R2', 'R-squared')
plot_metric('Time(s)', 'Time (s)')
```

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
```

```

# Sample fixed dummy data (replace with your full data)
data = {
    'Output': ['ST58TA', 'ST117TA', 'ST40TA', 'ST99TA', 'ST32TA', 'ST11TA', 'ST118TA', 'ST119TA'],
    'Model': ['RF + LGB + XGB'] * 10 + ['GB + XGB + CB'] * 10 + ['Ridge + LGB + XGB'] * 10,
    'MSE': [
        0.0076, 0.0065, 0.0092, 0.0081, 0.0095, 0.0086, 0.0224, 0.0076,
        0.0076, 0.0065, 0.0092, 0.0080, 0.0094, 0.0085, 0.0222, 0.0076,
        0.0047, 0.0042, 0.0066, 0.0052, 0.0064, 0.0057, 0.0142, 0.0047
    ],
    'RMSE': [
        0.0875, 0.0808, 0.0960, 0.0901, 0.0976, 0.0927, 0.1496, 0.0875,
        0.0875, 0.0808, 0.0960, 0.0896, 0.0971, 0.0923, 0.1491, 0.0875,
        0.0683, 0.0644, 0.0815, 0.0725, 0.0802, 0.0758, 0.1193, 0.0683
    ],
    'R2': [
        0.8334, 0.8517, 0.8084, 0.8233, 0.8085, 0.8171, 0.8904, 0.8334,
        0.8336, 0.8519, 0.8086, 0.8245, 0.8096, 0.8183, 0.8907, 0.8336,
        0.8970, 0.9031, 0.8563, 0.8924, 0.8711, 0.8895, 0.9228, 0.8970
    ],
    'Time(s)': [
        61.21, 57.84, 67.43, 63.67, 61.98, 60.92, 62.91, 59.86, 62.94,
        94.78, 90.52, 98.95, 96.55, 92.98, 91.45, 97.34, 90.23, 94.86,
        86.84, 83.23, 92.78, 90.97, 85.53, 82.54, 93.42, 86.42, 85.53
    ]
}

# Create DataFrame
df = pd.DataFrame(data)

# Plot settings for reuse
def plot_metric(metric, ylabel):
    plt.figure(figsize=(10, 6))
    for model in df['Model'].unique():
        subset = df[df['Model'] == model]
        plt.plot(subset['Output'], subset[metric], marker='o', line)
    plt.title(f'{ylabel} for each Model')
    plt.xlabel('Output')
    plt.ylabel(ylabel)
    plt.xticks(rotation=45)
    plt.legend()
    plt.tight_layout()
    plt.show()

# Plot each metric
plot_metric('MSE', 'MSE')
plot_metric('RMSE', 'RMSE')
plot_metric('R2', 'R-squared')
plot_metric('Time(s)', 'Time (s)')

```

```

In [ ]: import pandas as pd
import matplotlib.pyplot as plt

# Corrected data
data = {
    'Output': ['ST58TA', 'ST117TA', 'ST40TA', 'ST99TA', 'ST32TA', 'ST11TA', 'ST118TA', 'ST119TA'],

```

```

'Model Combination': ['RF + GB + XGB + LGB'] * 10 + ['RF + GB +
'RMSE': [
    0.0045, 0.0040, 0.0064, 0.0050, 0.0062, 0.0055, 0.0138, 0.0
    0.0044, 0.0040, 0.0064, 0.0050, 0.0062, 0.0055, 0.0138, 0.0
    0.0045, 0.0040, 0.0064, 0.0050, 0.0062, 0.0055, 0.0138, 0.0
    0.0043, 0.0039, 0.0062, 0.0049, 0.0061, 0.0054, 0.0135, 0.0
],
'R2': [
    0.8998, 0.9059, 0.8591, 0.8948, 0.8732, 0.8917, 0.9249, 0.9
    0.9005, 0.9063, 0.8597, 0.8952, 0.8737, 0.8921, 0.9252, 0.9
    0.8998, 0.9059, 0.8591, 0.8948, 0.8732, 0.8917, 0.9249, 0.9
    0.9022, 0.9080, 0.8615, 0.8967, 0.8756, 0.8936, 0.9272, 0.9
],
'Training Time (s)': [
    72.13, 68.95, 76.42, 73.67, 71.23, 70.11, 77.34, 69.24, 74.
    98.23, 94.65, 103.21, 99.87, 95.42, 92.54, 104.78, 96.35, 9
    86.14, 83.25, 89.42, 86.97, 82.23, 81.11, 89.34, 83.42, 84.
    92.24, 89.34, 95.42, 92.57, 88.24, 87.11, 96.34, 88.67, 90.
]
}

df = pd.DataFrame(data)

# Plot for RMSE
plt.figure(figsize=(10, 6))
for model in df['Model Combination'].unique():
    subset = df[df['Model Combination'] == model]
    plt.plot(subset['Output'], subset['RMSE'], marker='o', linestyle=
plt.title('RMSE for each Model Combination')
plt.xlabel('Output')
plt.ylabel('RMSE')
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()

# Plot for R2
plt.figure(figsize=(10, 6))
for model in df['Model Combination'].unique():
    subset = df[df['Model Combination'] == model]
    plt.plot(subset['Output'], subset['R2'], marker='o', linestyle=
plt.title('R-squared for each Model Combination')
plt.xlabel('Output')
plt.ylabel('R-squared')
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()

# Plot for Training Time
plt.figure(figsize=(10, 6))
for model in df['Model Combination'].unique():
    subset = df[df['Model Combination'] == model]
    plt.plot(subset['Output'], subset['Training Time (s)'], marker=
plt.title('Training Time for each Model Combination')
plt.xlabel('Output')

```

```
plt.ylabel('Training Time (s)')
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()
```

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt

# Sample complete data (please replace `...` with actual values)
data = {
    'Output': ['ST58TA', 'ST117TA', 'ST40TA', 'ST99TA', 'ST32TA', 'ST11TA', 'ST118TA', 'ST119TA', 'ST120TA', 'ST121TA'],
    'Model Combination': ['RF + GB + XGB + LGB + CB'] * 10,
    'RMSE': [0.0043, 0.0039, 0.0062, 0.0049, 0.0061, 0.0054, 0.0135, 0.0135, 0.0135, 0.0135],
    'R2': [0.9025, 0.9083, 0.8618, 0.8970, 0.8759, 0.8939, 0.9275, 0.9275, 0.9275, 0.9275],
    'Training Time (s)': [110.45, 107.68, 115.23, 111.87, 106.23, 106.23, 106.23, 106.23, 106.23, 106.23]
}

df = pd.DataFrame(data)

# Plot for RMSE
plt.figure(figsize=(10, 6))
plt.plot(df['Output'], df['RMSE'], marker='o', linestyle='-', label='RMSE')
plt.title('RMSE for each Model Combination')
plt.xlabel('Output')
plt.ylabel('RMSE')
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()

# Plot for R-squared
plt.figure(figsize=(10, 6))
plt.plot(df['Output'], df['R2'], marker='o', linestyle='-', label='R-squared')
plt.title('R-squared for each Model Combination')
plt.xlabel('Output')
plt.ylabel('R-squared')
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()

# Plot for Training Time
plt.figure(figsize=(10, 6))
plt.plot(df['Output'], df['Training Time (s)'], marker='o', linestyle='-', label='Training Time (s)')
plt.title('Training Time for each Model Combination')
plt.xlabel('Output')
plt.ylabel('Training Time (s)')
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()
```

```
In [ ]:
```